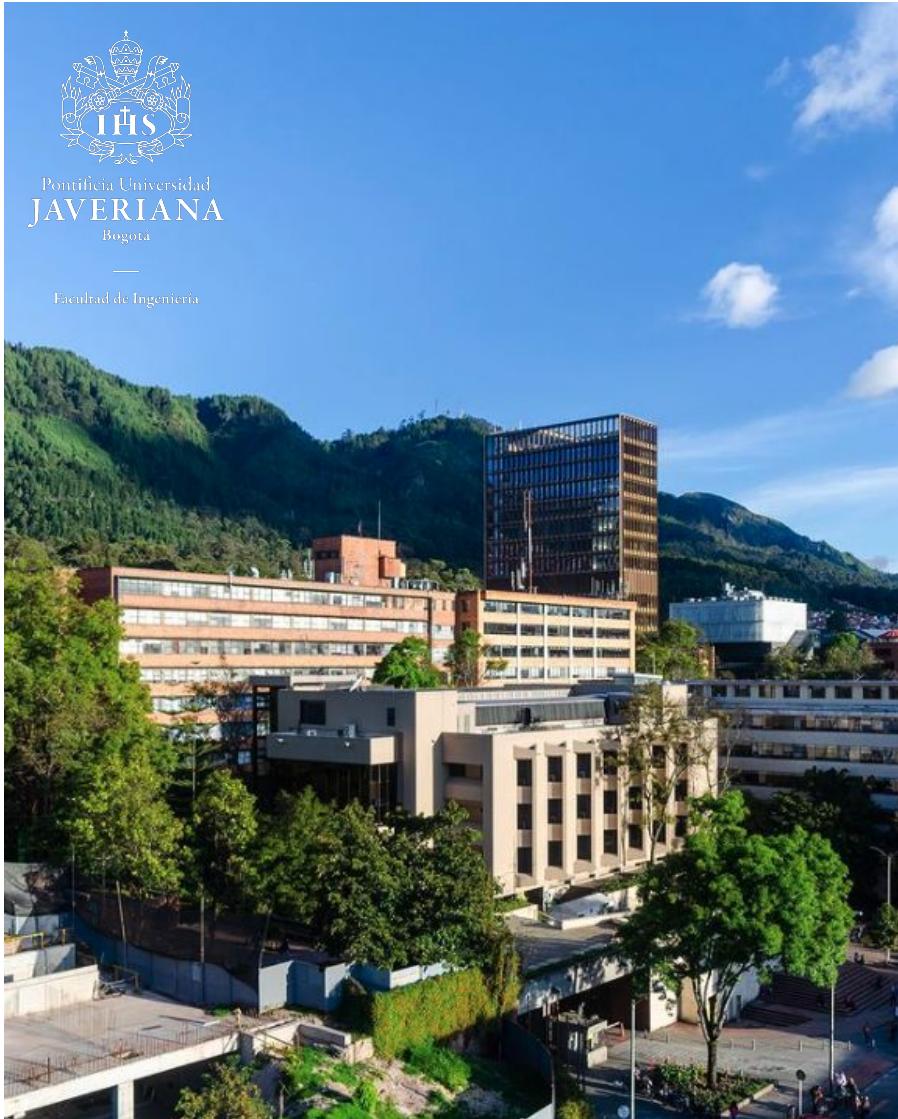




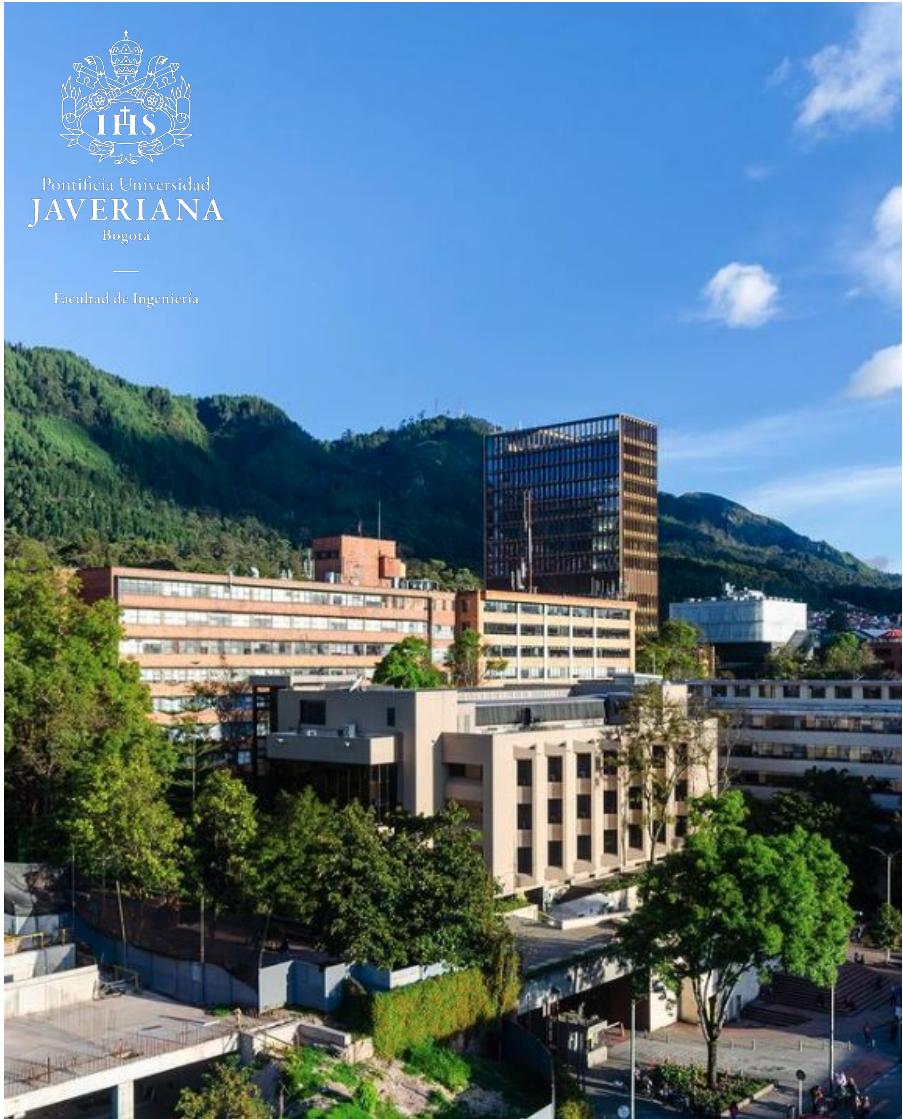
Machine Learning Fundamentals

Workshop para América Latina y el Caribe (WALC)
Track 2 – Inteligencia Artificial Aplicada
November 14, 2023

Pontificia Universidad Javeriana – Bogotá, Colombia



Pontificia Universidad Javeriana – Bogotá, Colombia

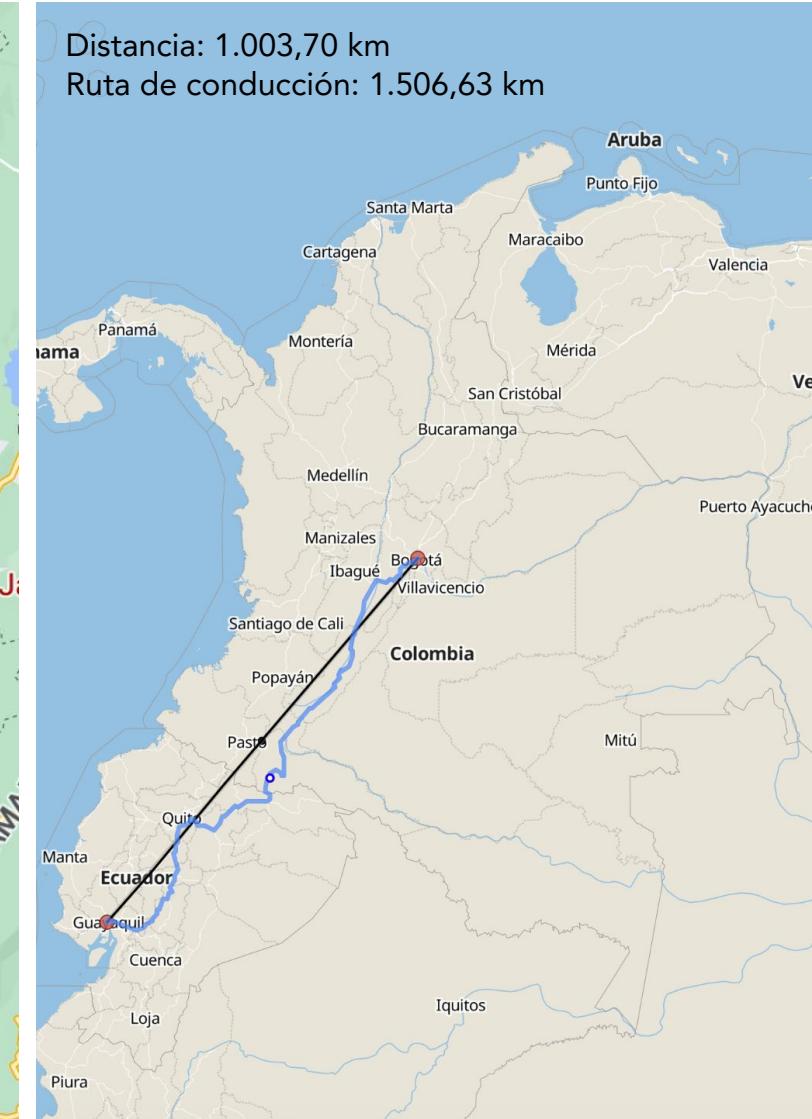


Pontificia Universidad
JAVERIANA
Bogotá

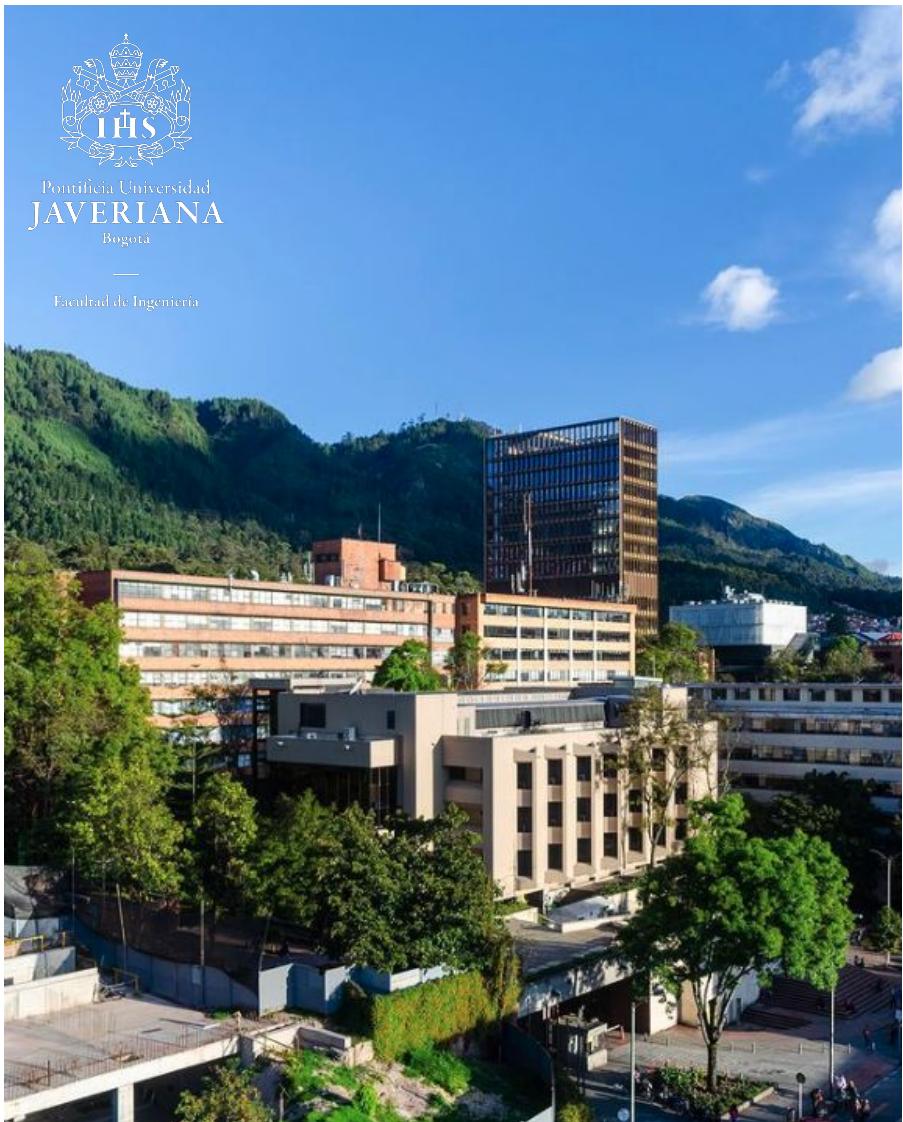
Facultad de Ingeniería



Distancia: 1.003,70 km
Ruta de conducción: 1.506,63 km

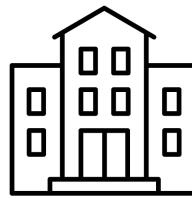


Pontificia Universidad Javeriana



Pontificia Universidad
JAVERIANA
Bogotá

Facultad de Ingeniería



18 Schools



62 Departments



15 Institutes



32.000

Students

42 >

Undergraduate



27.000

Undergraduate
Students

92 >

Specializations



1.700

Employees

67 >

Masters

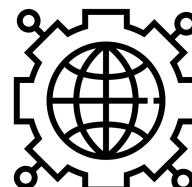


160.000

Alumni

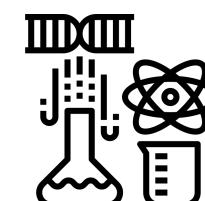
15 >

PhD. programs



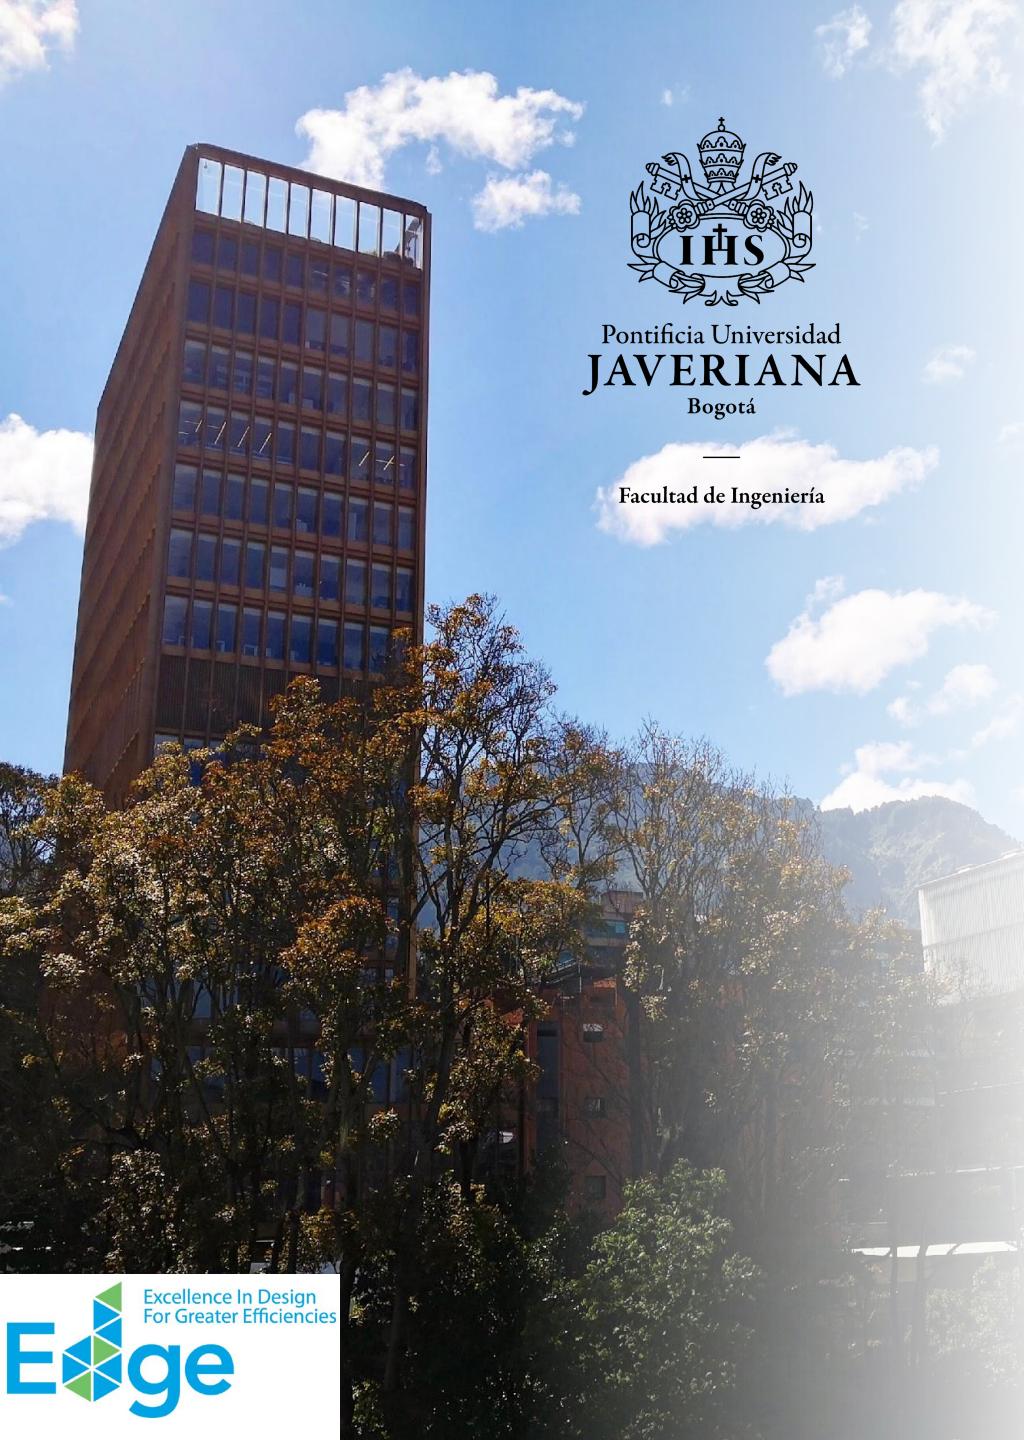
5

Centers of Excellence



96

Research Groups



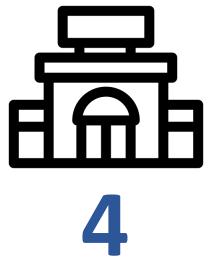
School of Engineering

- 3.750 students (3000 undergrad)
- 104 professors
- 20.000 alumni

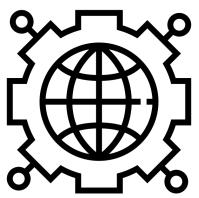
- 8 undergraduate programs
- 7 specializations
- 12 master programs
- 2 doctoral programs

New Research and Laboratories Building

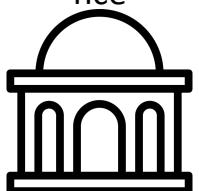
- 14.082 m² (700 for student workspaces)
- 93 meters high
- 15 floors and 3 basements
- 10.000 pieces of state-of-the-art equipment
- <https://ingenieria.javeriana.edu.co/nuestro-edificio>



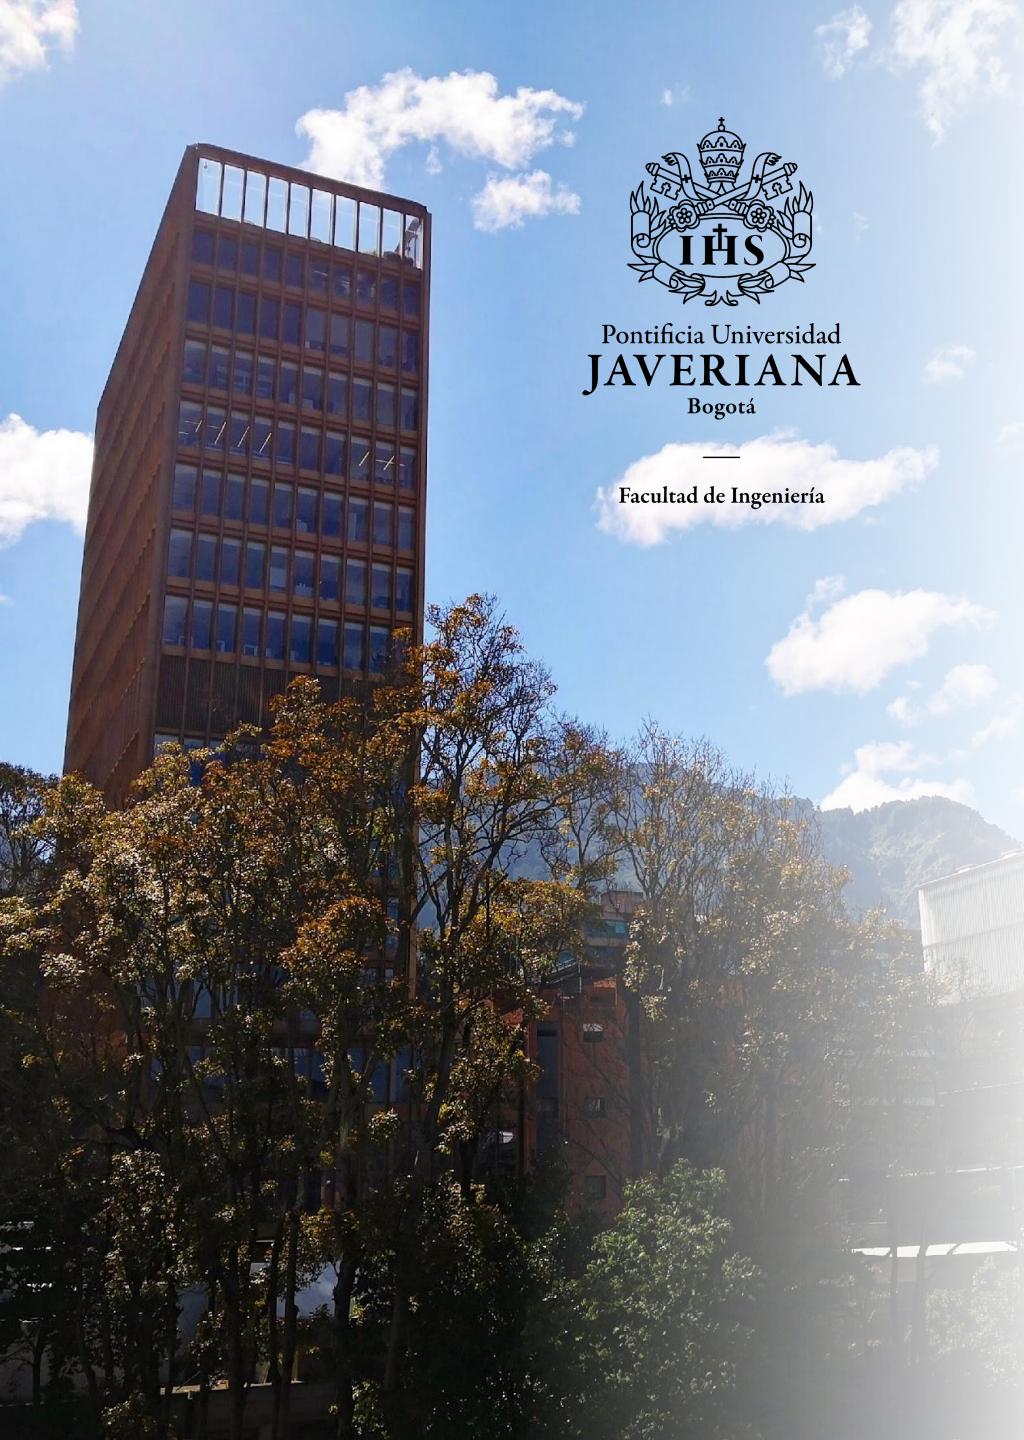
Departments



Centers of Excellence



Institutes



Our Programs

Undergraduate

- Electronics Eng.
- Civil Eng.
- Industrial Eng.
- Systems and Computers Eng.
- Mechatronics Eng.
- Telecommunications and Networks Eng.
- Bioengineering
- Mechanical Eng.
- Data Science



Graduate

- Doctorate in Engineering
- Doctorate in Science and Technology of Materials
- Master in Electronics Eng.
- Master in Civil Eng.
- Master in Hydro-systems Eng.
- Master in Systems and Computing Eng.
- Master in Bioengineering
- Master in Analytics
- Master in Industrial Eng.
- Master in Logistics and Transport
- Master in Energy and Sustainability
- Master in Digital Security **NEW**
- Master in Internet of Things **NEW**
- Master in Artificial Intelligence **NEW**

Diego Méndez Chaves, Ph.D

School of Engineering – Pontificia Universidad Javeriana

2009-2
012 Ph.D and M.Sc in
 Computer Science
 Full Scholarship

University of South Florida
(USA)

Participatory sensing and
location-based systems

2005-2
008 M.Eng in Electronics
 Engineering and
 Computers
 Full Scholarship

Universidad de Los Andes
(Colombia)

Wireless sensor networks

2000-
2004 B.Eng in Electronics
 Engineering

Universidad Nacional de
Colombia

Embedded systems design



Diego Méndez Chaves, Ph.D

School of Engineering – Pontificia Universidad Javeriana

Universidad
Nacional de Colombia

2006: Instructor Professor
Teaching activities

Universidad de
Los Andes

2005-2009: Graduate Assistant and Instructor Professor
Teaching and research activities

University of
South Florida

2009-2012: Graduate Assistant and Instructor Professor
Teaching and research activities

Pontificia
Universidad Javeriana

2012 - to date: Associate Professor
Teaching, research and service activities

External Collaborations

2016: Invited Researcher at the IHP Microelectronics
(Frankfurt-Oder, Germany)
2017-2018: Advising Engineering Director and Project Manager
in IoT (IKUSI, Colombia)

- Teaching: 21 undergraduate (7), 26 master (7), 5 doctoral (3).
- Research: ~3 international cooperation (2), 1 center of excellence, >10 industry related, >20 research, ~59 publications.





Diego Méndez Chaves, Ph.D

School of Engineering – Pontificia Universidad Javeriana

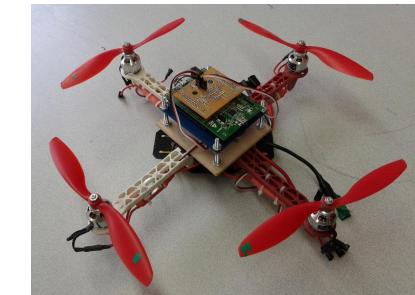
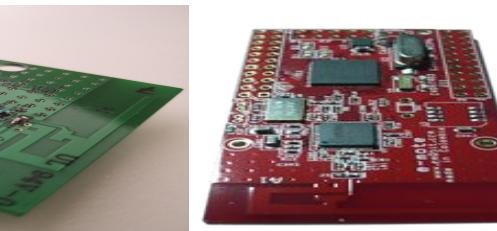
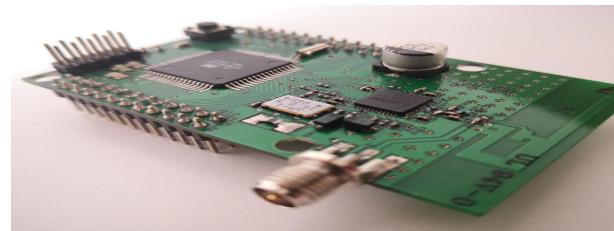
- Associate Professor at the Department of Electronics Engineering.
- Director of the Master Program in Internet of Things.
- Director of the Master Program in Electronics Engineering.
- Technical Director of the Center of Excellence and Adoption in IoT (CEA-IoT)
- Research associate at the Marconi Lab in the International Centre for Theoretical Physics (ICTP), Trieste - Italy.
- TinyML Academic Network – Founding member and local coordinator.
- Research interests: IoT, embedded systems, wireless sensor networks, participatory sensing, digital systems design and embedded operating systems.

Embedded Research

Digital oscilloscope and a logic analyzer on a GBA



DA-MOTE: A Flexible Platform for WSN Apps

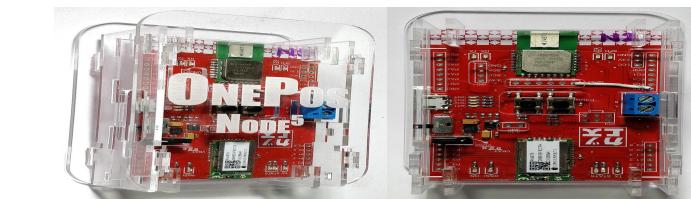


A completely customized UAV for research projects

Frictionless Indoor Location using BLE



OpenWuR: Open Platform for WuR-based Apps



4D Phenotyping with LiDAR and Multispectral Images



The one
that
started
all!

Embedded Teaching

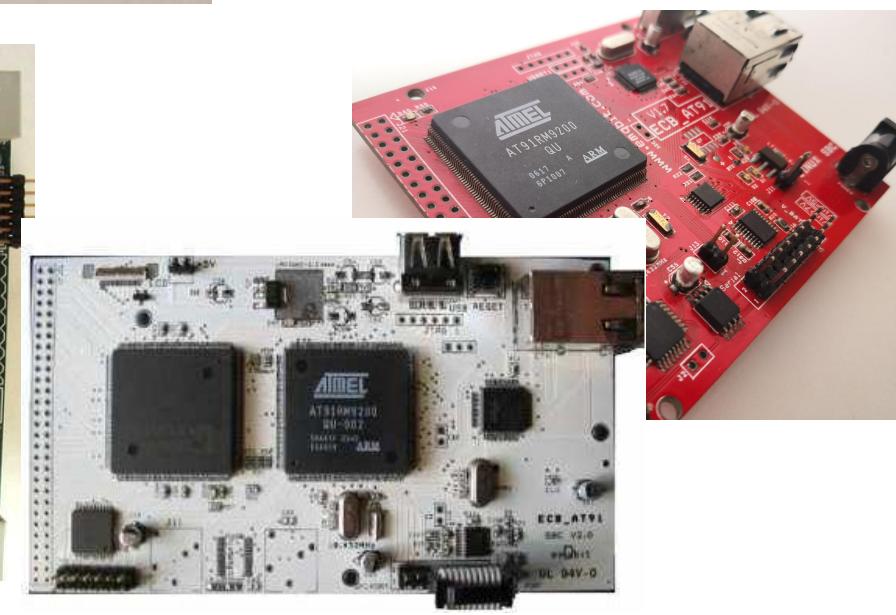
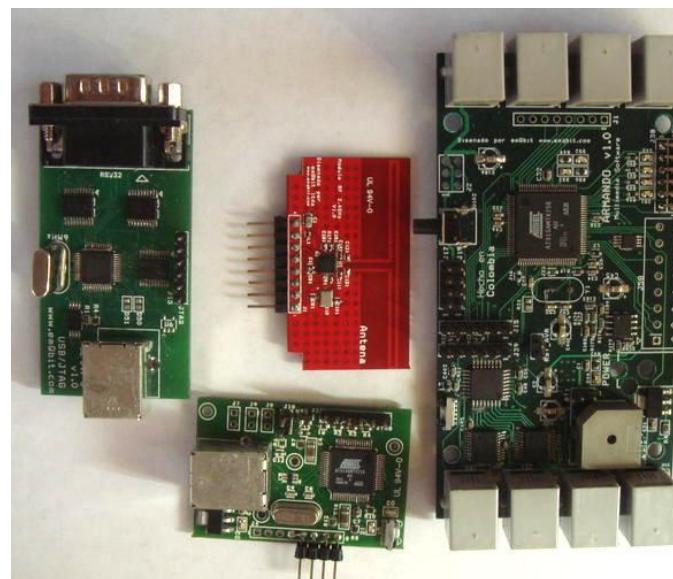
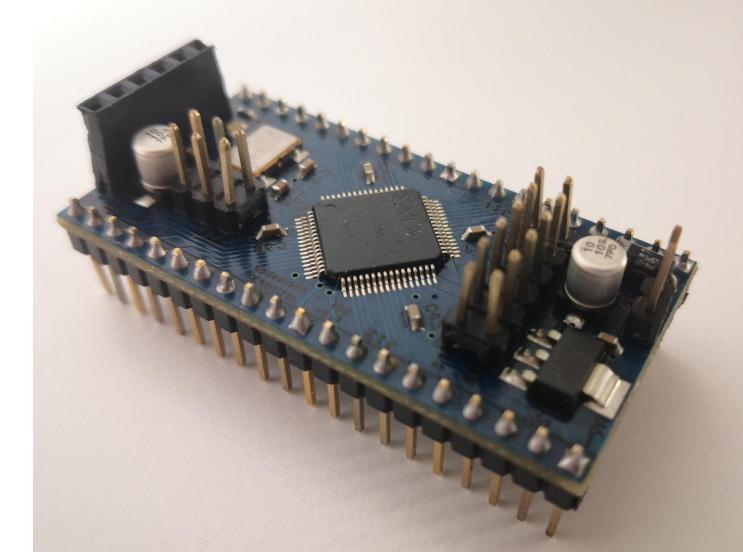
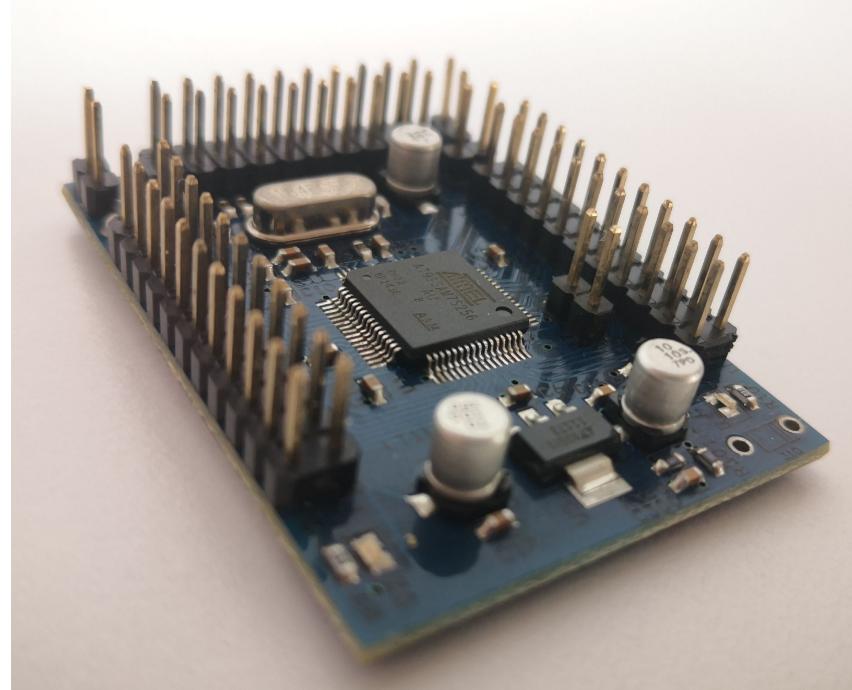
Design with
Microcontrollers (SAM7S)

Digital Hardware Design
(XC9572XL)

STEM-oriented robotics
platform

Embedded Linux projects
(AT91RM9200)

Co-design projects
(SoC + FPGA)



Embedded Service

Agrosavia



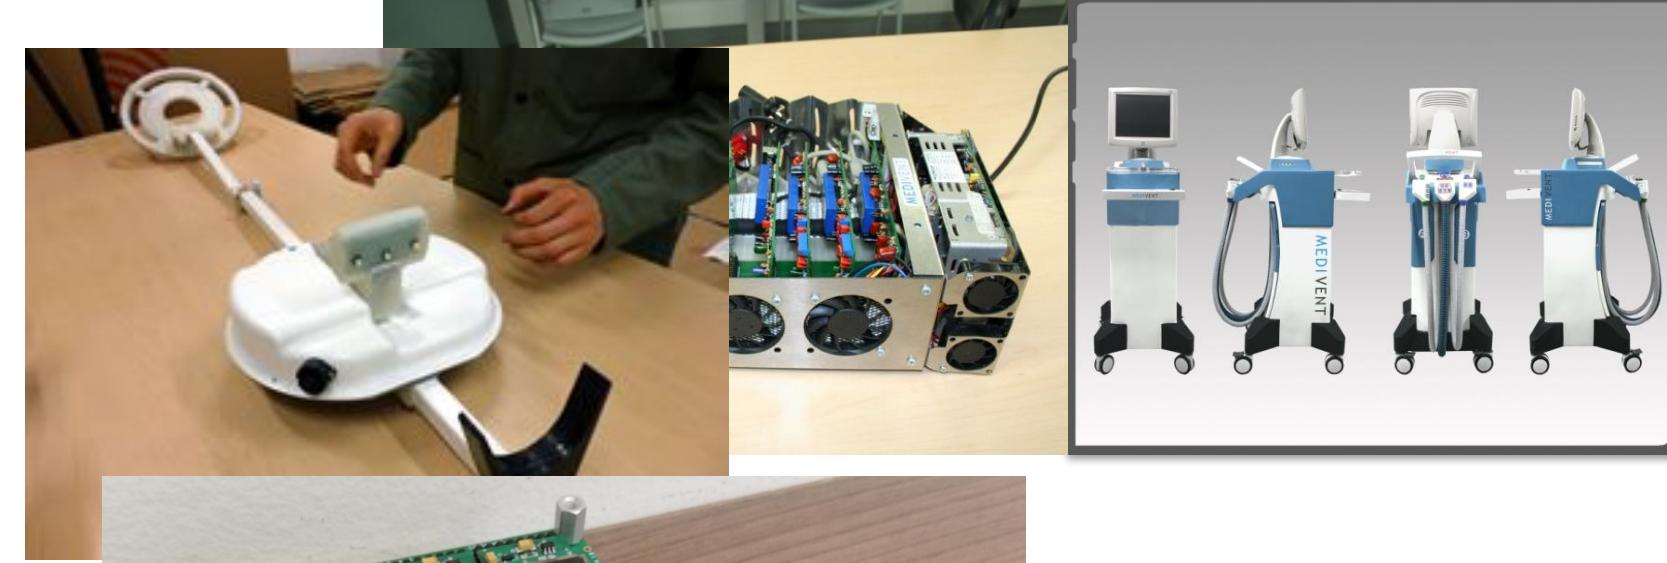
Cornare

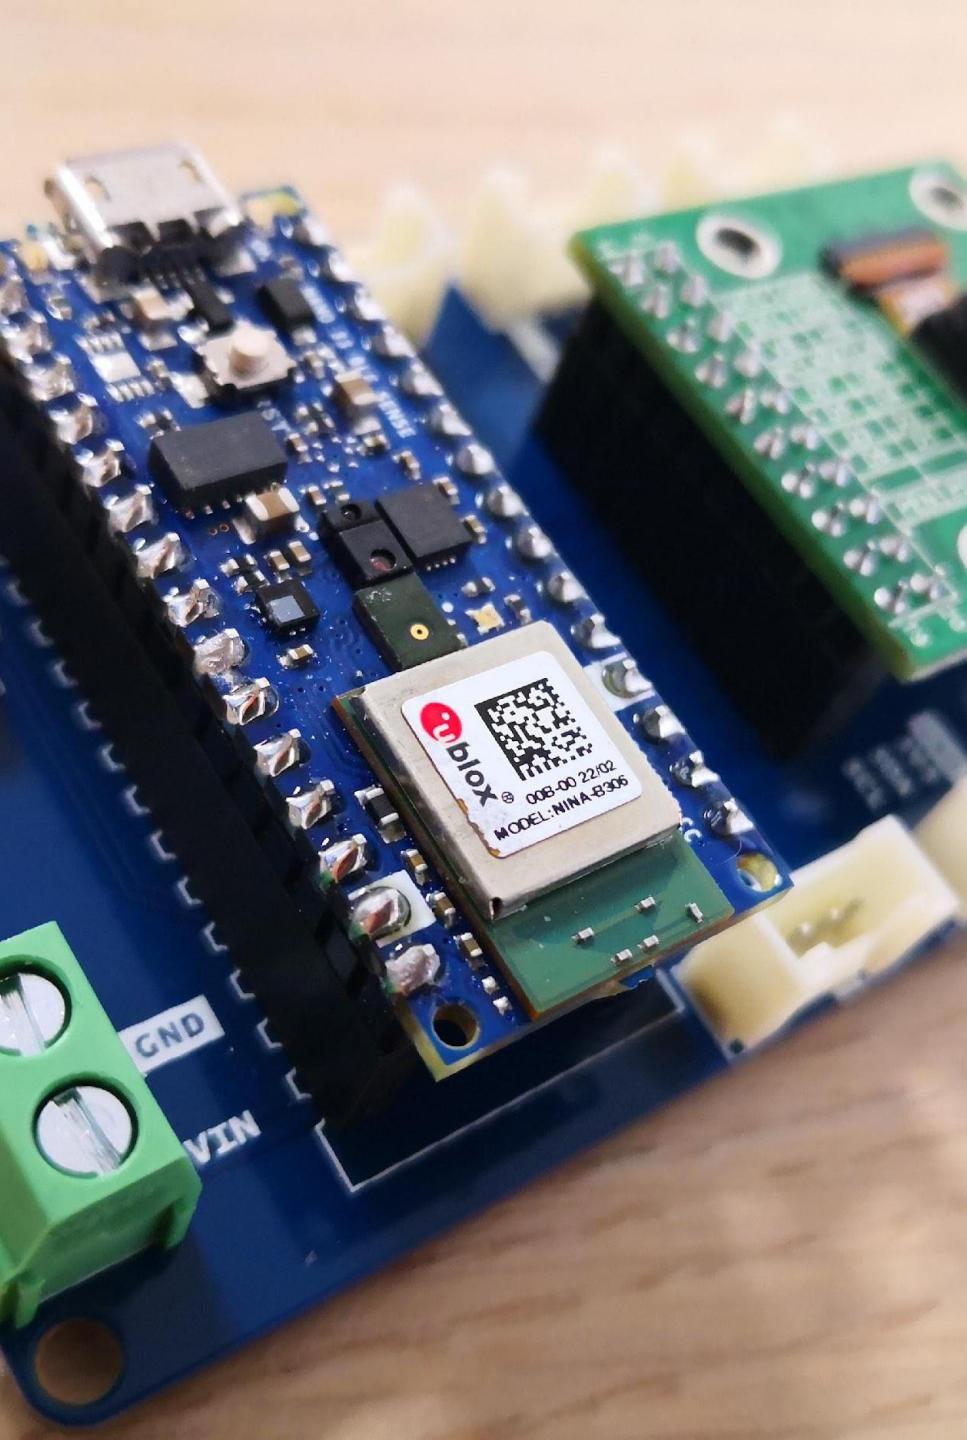
PIMA

IKUSI

Indumil

Medivent





Outline

- AI vs ML vs DL
- TinyML's Main Components
- The Machine Learning Paradigm
- Finding the Best Solution and Fitting a Model
- Regression and Classification with NN
- ML Issues
- Challenges of TinyML





Alexa, play some
rock music!

Playing the Rock
Hits playlist

ElephantEdge

Building The World's Most Advanced **Wildlife Tracker**.



Source: <https://wildlabs.net/resources/competition/challenge-elephantedge>



ElephantEdge

Risk Monitoring

"Know when an elephant is moving into a high-risk area and send real-time notifications to park rangers."

Conflict Monitoring

"Sense and alert when an elephant is heading into an area where farmers live."

Activity Monitoring

"Classify the general behavior of the elephant, such as when it is drinking, eating, sleeping, etc."

Communication Monitoring

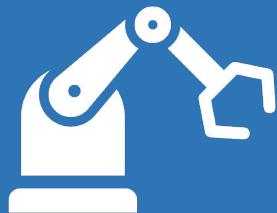
"Listen for vocal communications between elephants via the onboard microphone."

AI vs. ML vs. DL

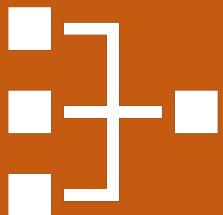
Artificial Intelligence



Machine Learning



Deep Learning



Artificial intelligence (AI): any technique that enables computers to mimic human intelligence.

Machine learning (ML): a subset of AI that uses techniques that enable machines to use experience to improve at tasks.

Deep learning (DL): a subset of ML based on artificial neural networks (ANN). The learning process is deep because of its structure.

General Steps for Machine Learning

On a high level, the craft of creating machine learning (ML) processes is comprised of several steps:

Decide on the Question

Collect and Prepare Data

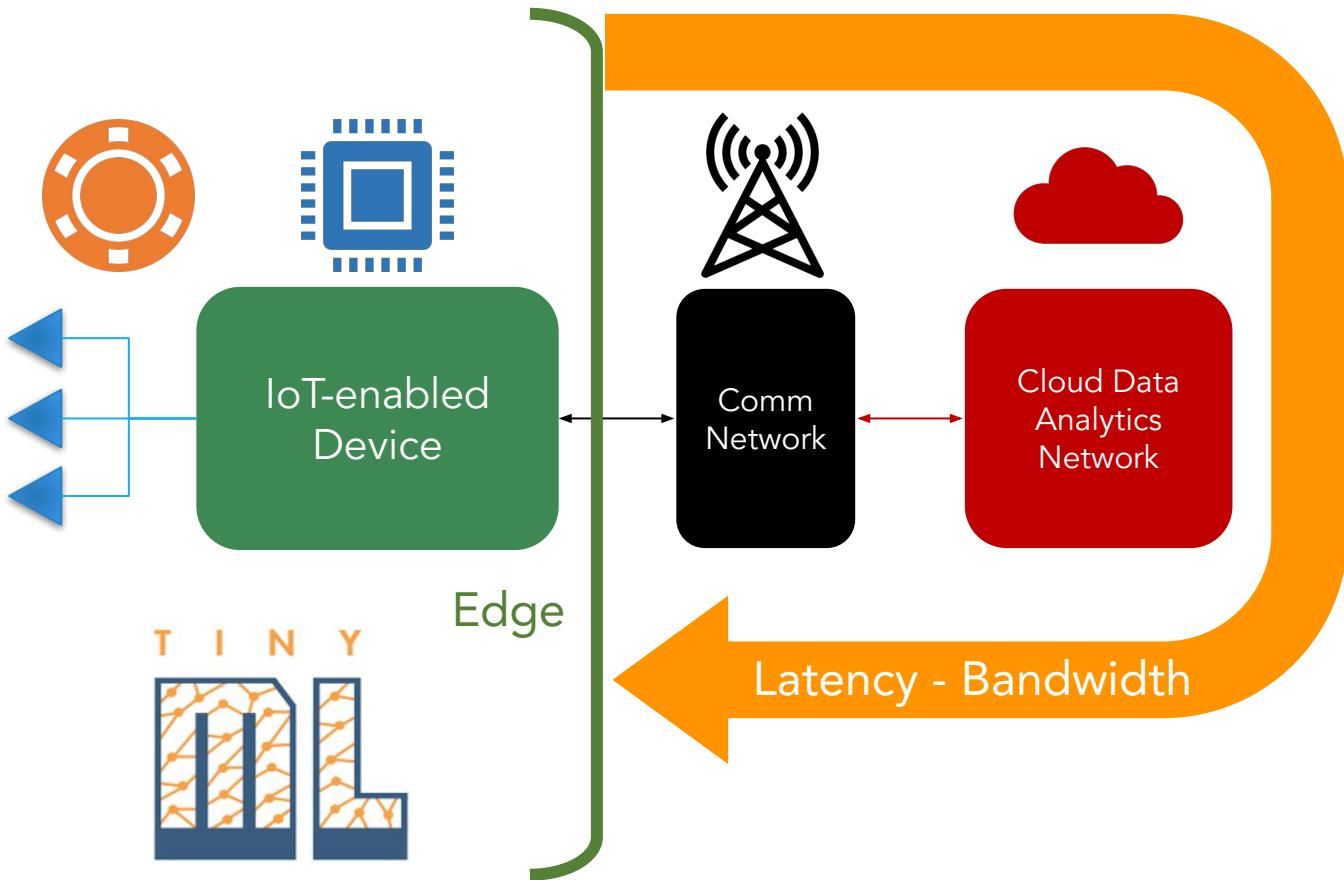
Choose a Training Method

Train the Model

Evaluate the Model

Parameter Tuning

Predict



“The future of ML is *tiny* and
bright.”

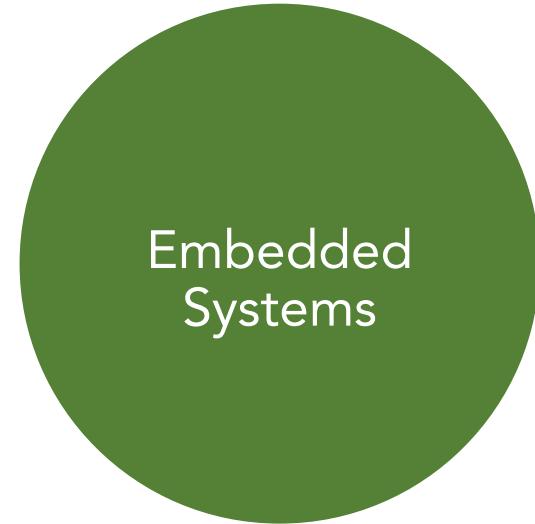
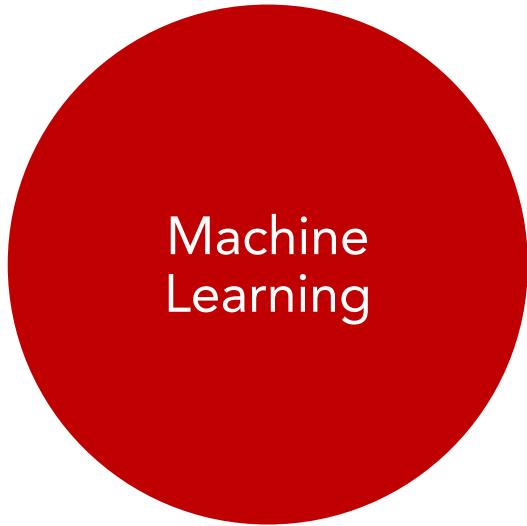
3 main components

"Edge AI is a truly complete technology. As a topic, it makes use of knowledge from everything from the physical properties of semiconductor electronics all the way up to the engineering of high-level architectures that span devices and the cloud. It demands expertise in the most cutting-edge approaches to **artificial intelligence and machine learning** along with the most venerable skills of bare-metal **embedded software engineering**. It makes use of the entire history of computer science and electrical engineering, laid out end to end."



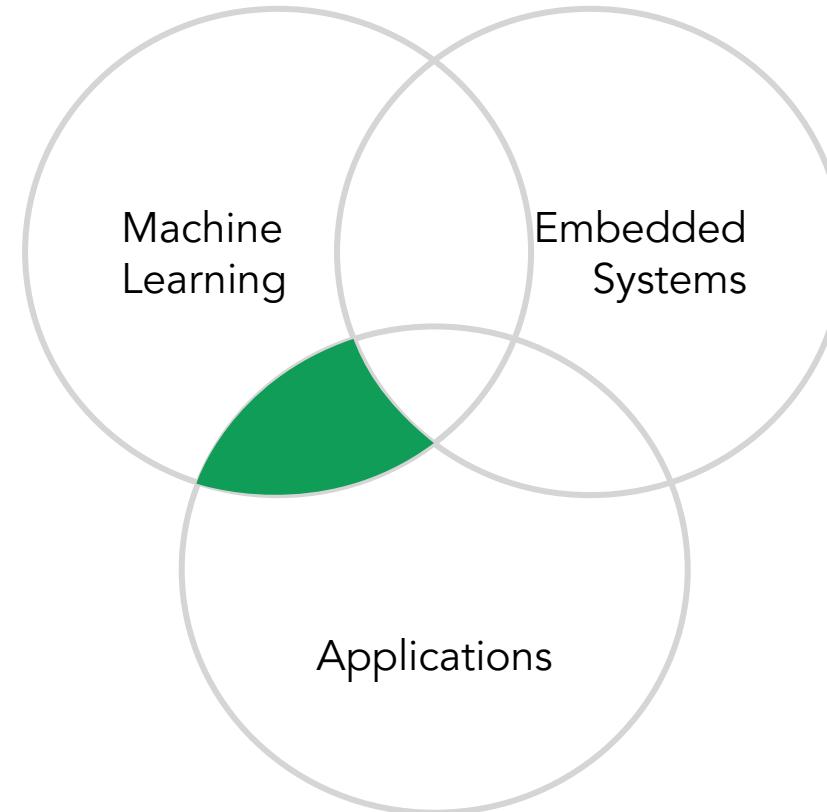
Situnayake, Daniel; Plunkett, Jenny
AI at the Edge (pp. 215-216)
O'Reilly Media

3 main components



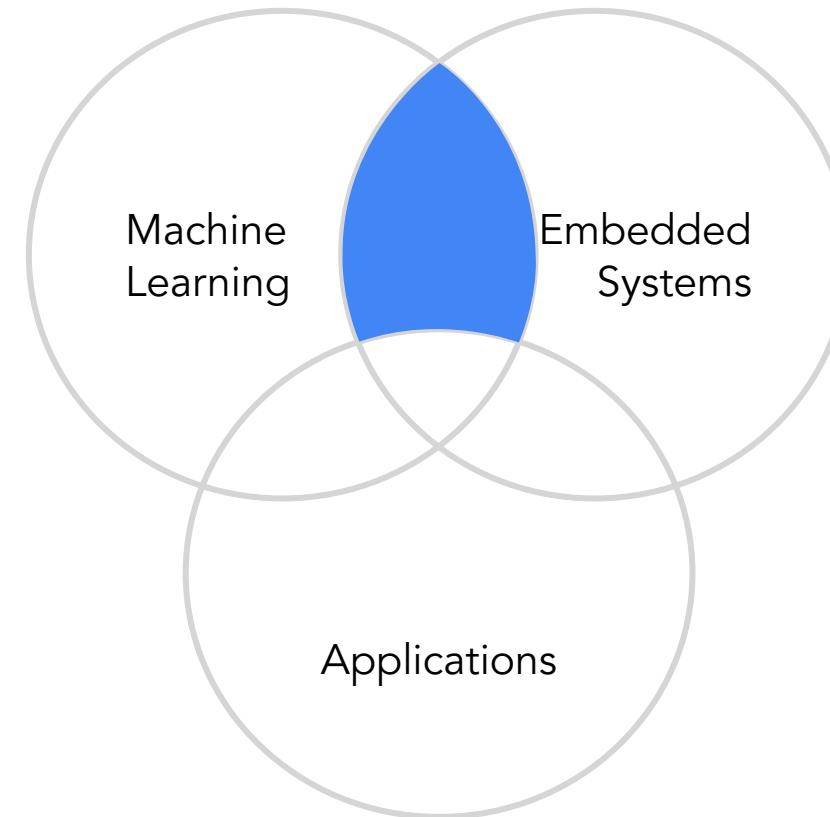
Interactions

How **machine learning** can
enable new and interesting
TinyML applications?



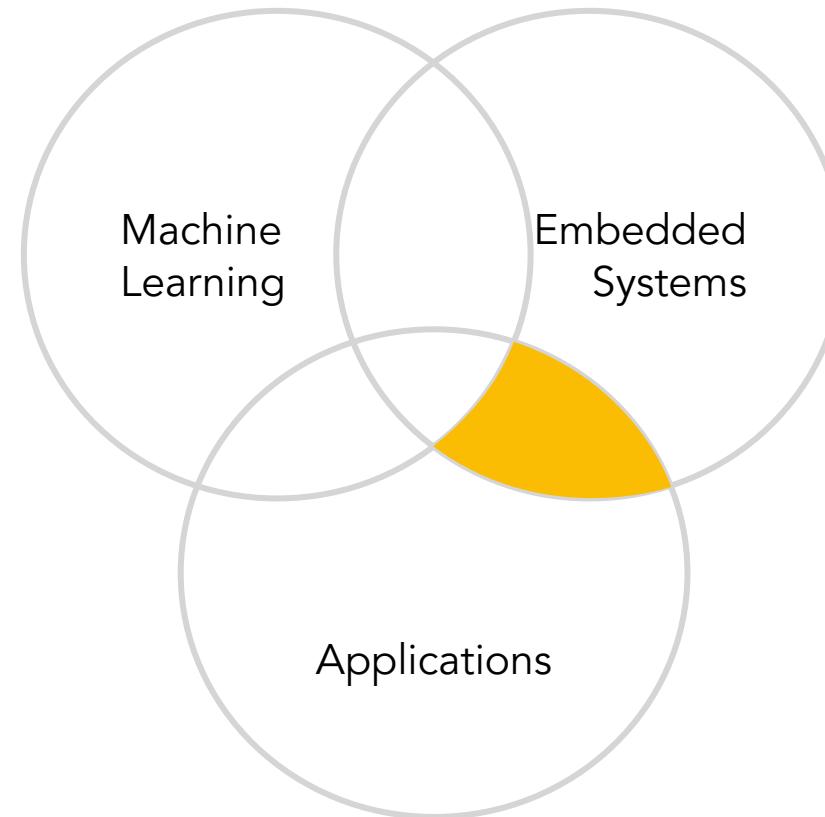
Interactions

What are the **challenges** with enabling **machine learning** on **tiny**, resource-constrained **embedded devices**?



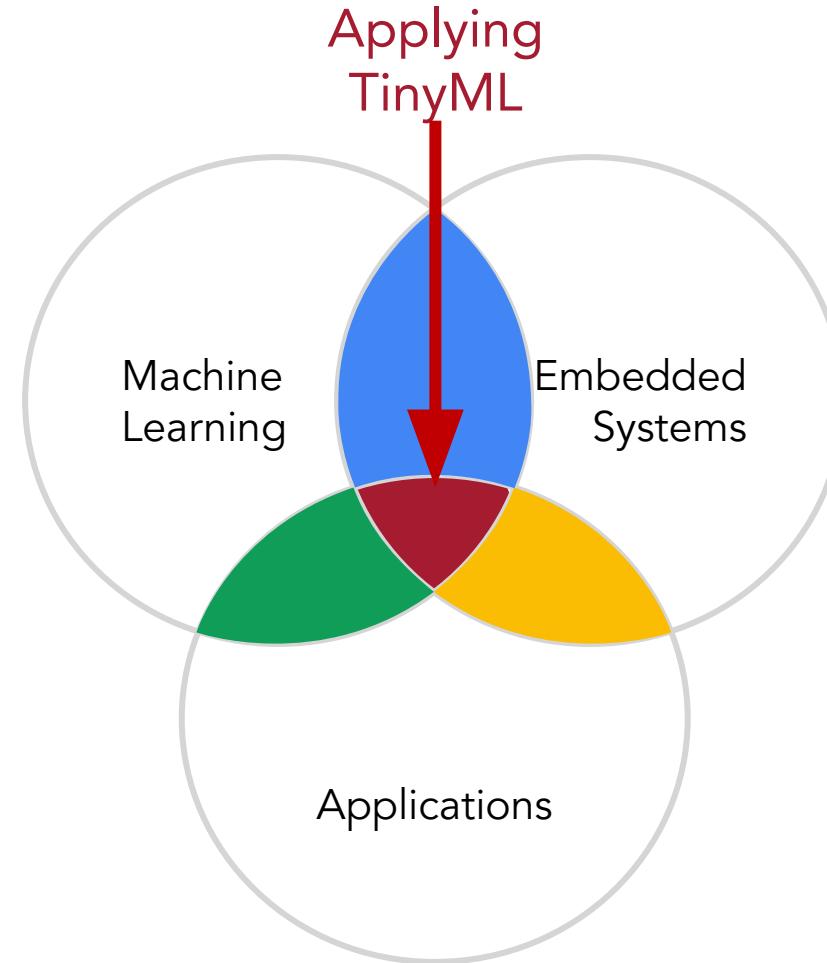
Interactions

What type of new **use cases** can we possibly enable on **embedded systems** that we could not otherwise do before?



Interactions

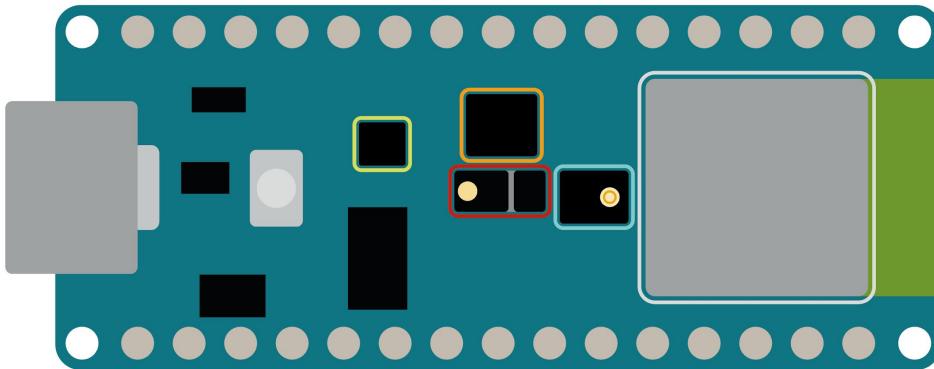
Given your understanding of things at these various intersections, you will have a deep understanding for **how to apply TinyML**



Our platform for experimental deployments

Arduino Nano 33 BLE Sense

NANO 33 BLE SENSE



- ◆ Color, brightness, proximity and gesture sensor
- ◆ Digital microphone
- ◆ Motion, vibration and orientation sensor
- ◆ Temperature, humidity and pressure sensor
- ◆ Arm Cortex-M4 microcontroller and BLE module



We will run through this long process



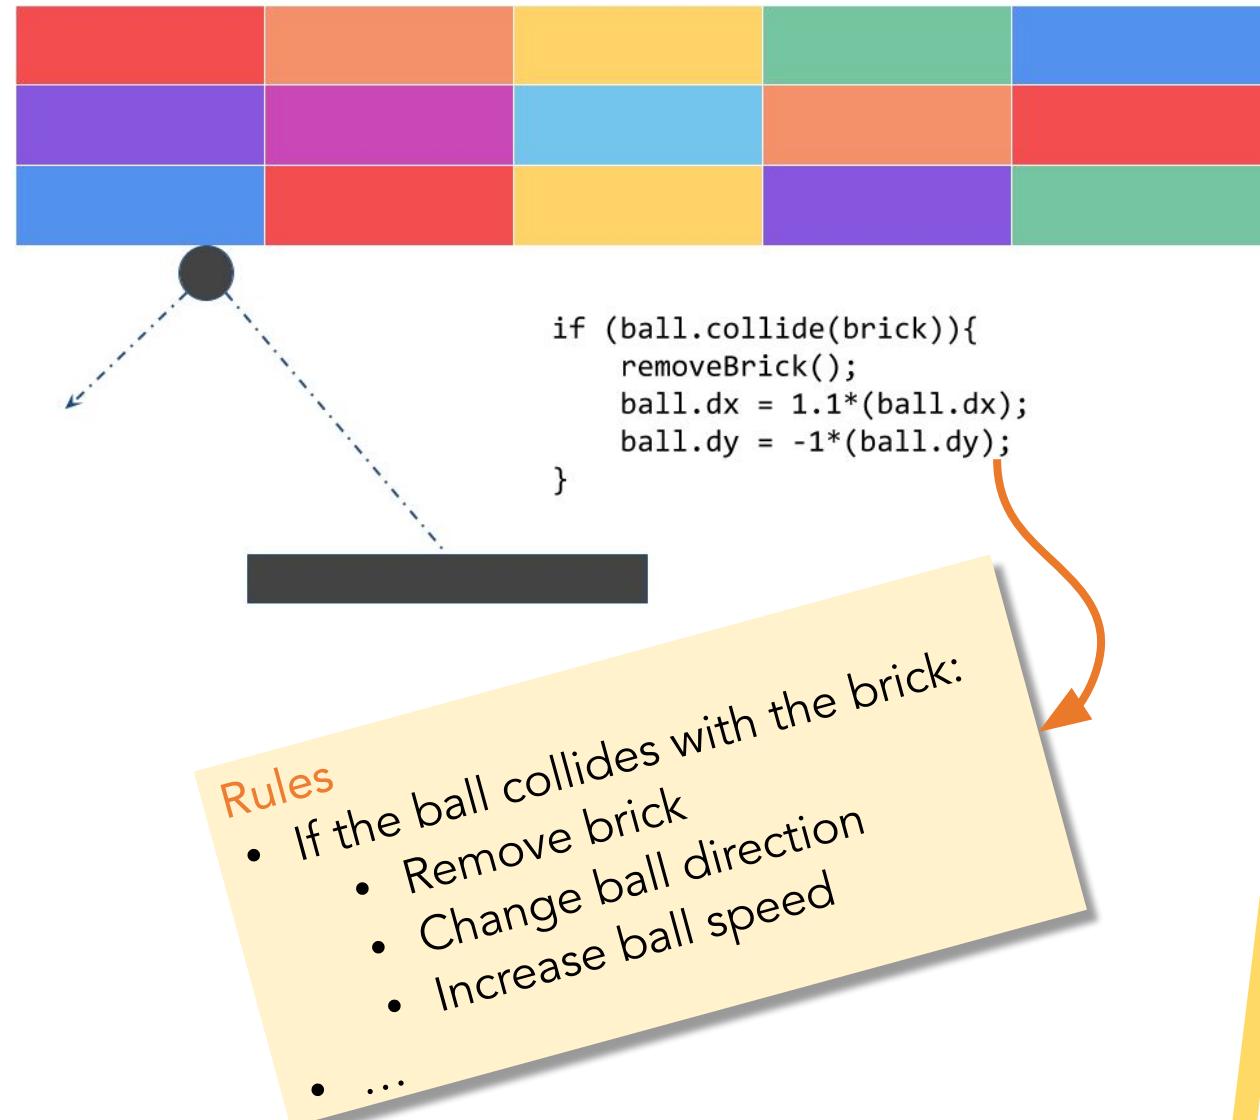
This is a **first encounter with ML**, but many things will be left to be **experimented or developed**.



The Machine Learning Paradigm

Explicit Coding

- Defining rules that determine behavior of a program
- Everything is pre-calculated and pre-determined by the programmer
- Scenarios are limited by program complexity



[https://en.wikipedia.org/wiki/Breakout_\(video_game\)](https://en.wikipedia.org/wiki/Breakout_(video_game))

The Traditional Programming Paradigm



Consider Activity Detection



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



// ???

Way too complex to code!

The Traditional Programming Paradigm



The Machine Learning Paradigm



Activity Detection with Machine Learning



0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010

1010100101001010101
0101010010010010001
001001111010101111
1010100100111101011

1001010011111010101
1101010111010101110
1010101111010101011
1111110001111010101

111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110

Label = WALKING

Label = RUNNING

Label = BIKING

Label = GOLFING

The Machine Learning Paradigm



0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010

Label = WALKING

1010100101001010101
0101010010010010001
0010011111010101111
1010100100111101011

Label = RUNNING

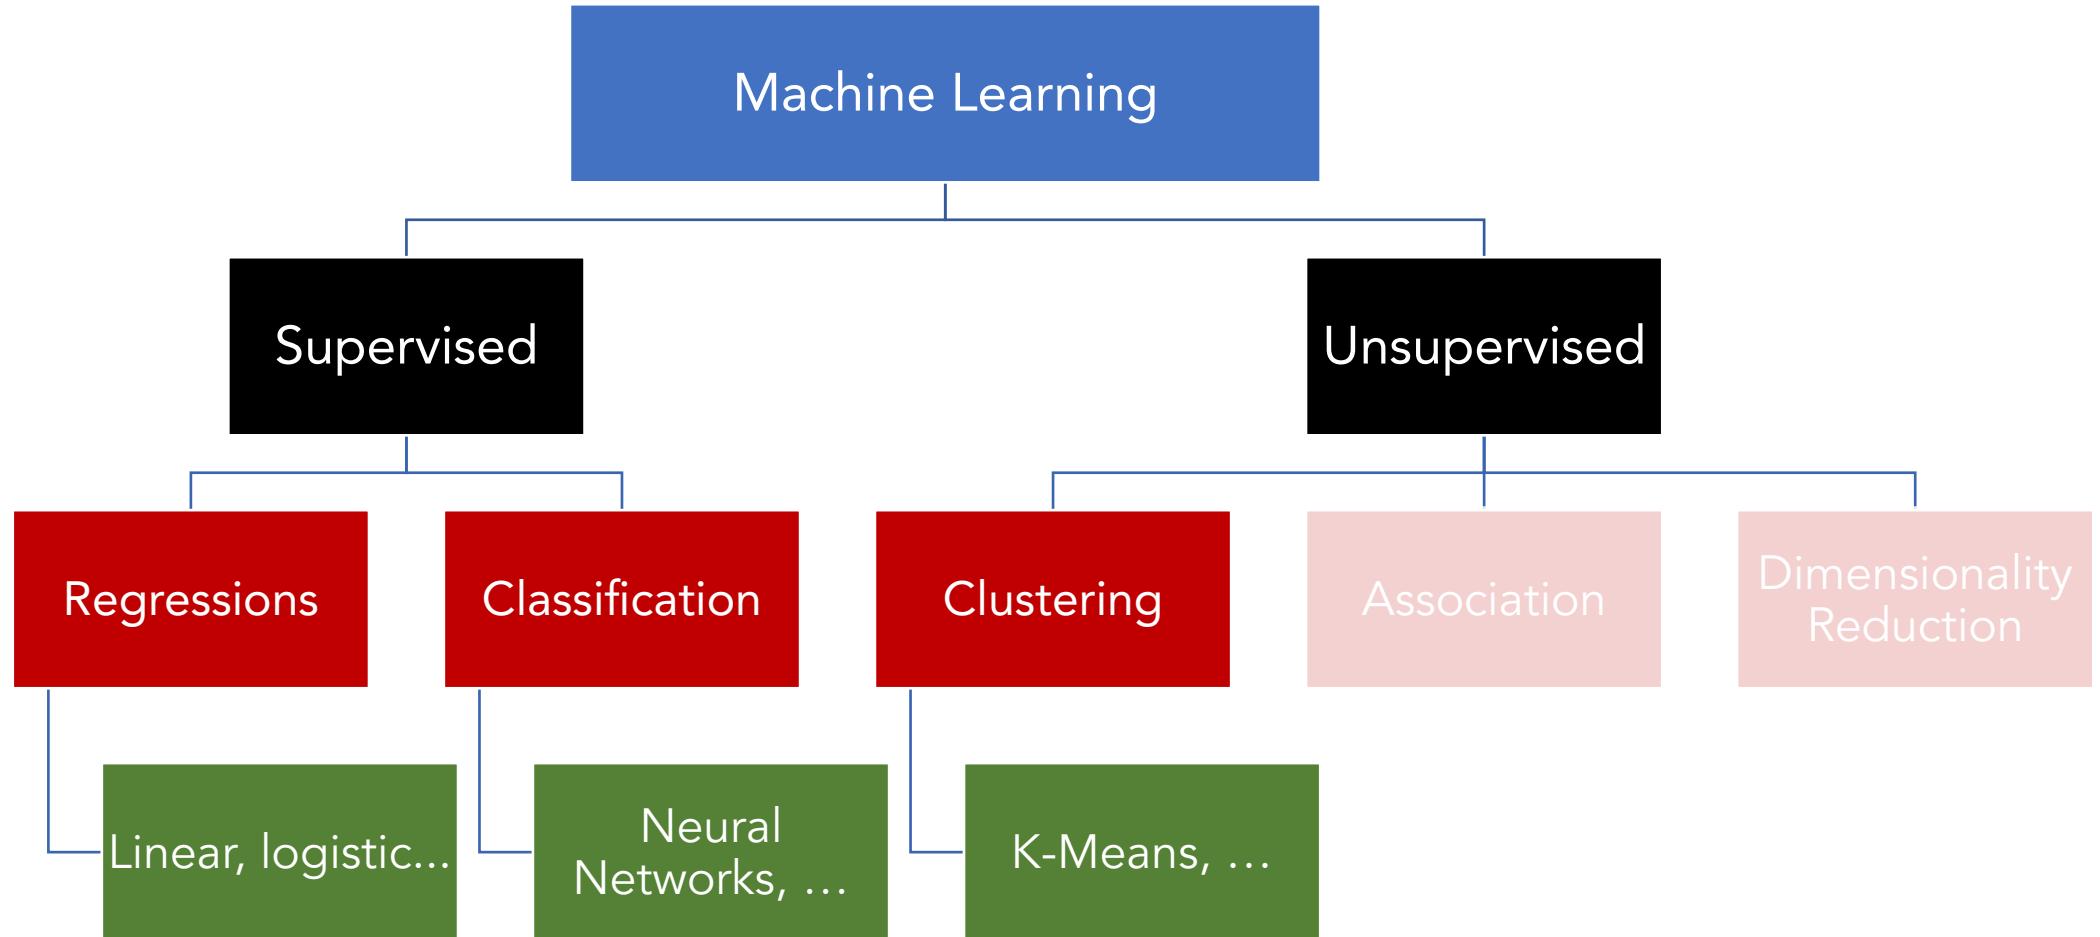
1001010011111010101
1101010111010101110
1010101111010101011
1111110001111010101

Label = BIKING

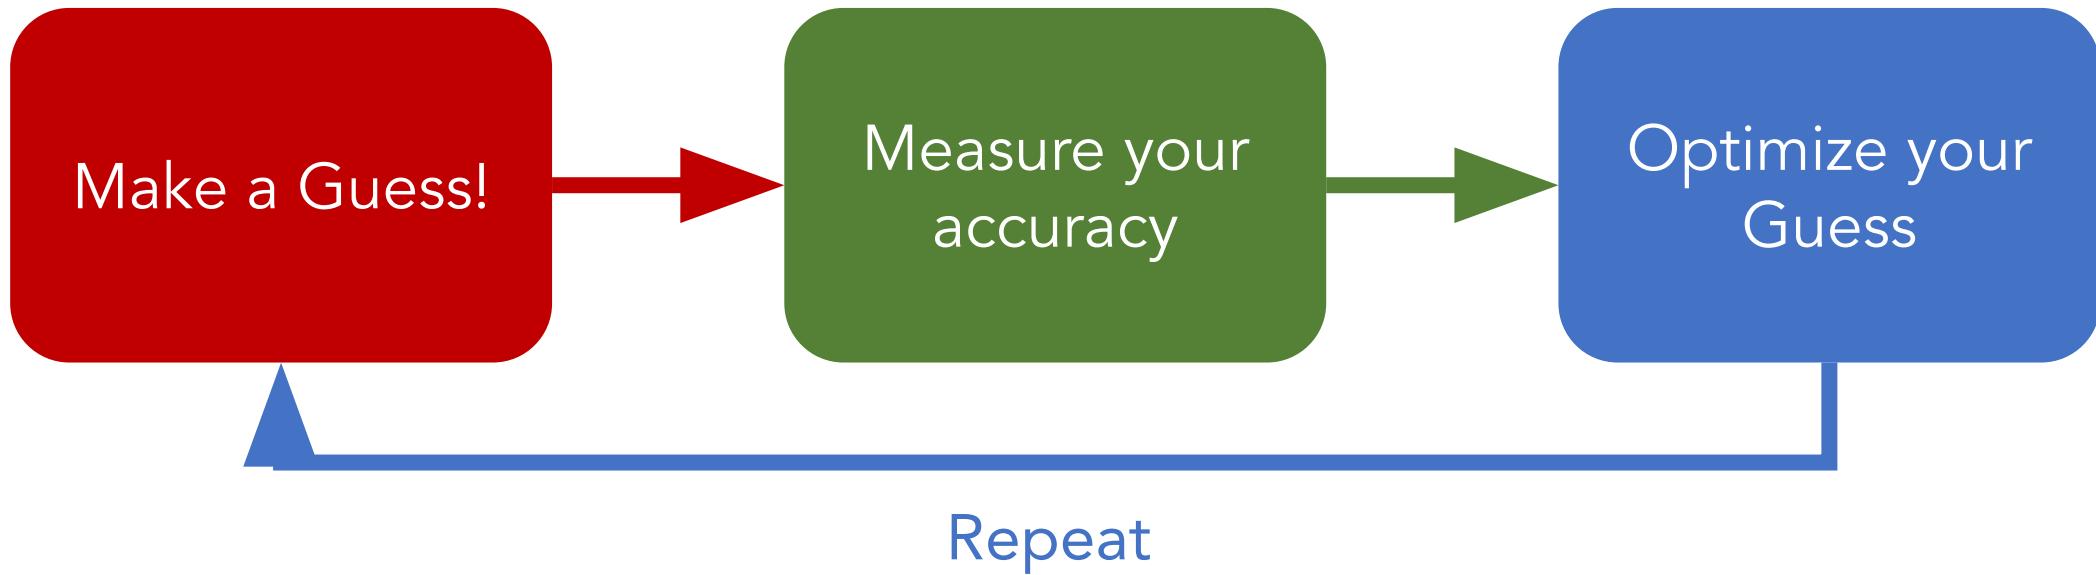
111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110

Label = GOLFING

Two Approaches



The Machine Learning Paradigm



The Machine Learning Paradigm



The Machine Learning Paradigm



How good is your model?

a way to measure your accuracy

Matching X to Y

X = { -1, 0, 1, 2, 3, 4 }

Y = { -3, -1, 1, 3, 5, 7 }



Make a guess!

$$Y = 3X - 1$$

$$X = \{ -1, 0, 1, 2, 3, 4 \}$$

$$\text{My } Y = \{ -4, -1, 2, 5, 8, 11 \}$$

How good is the guess?

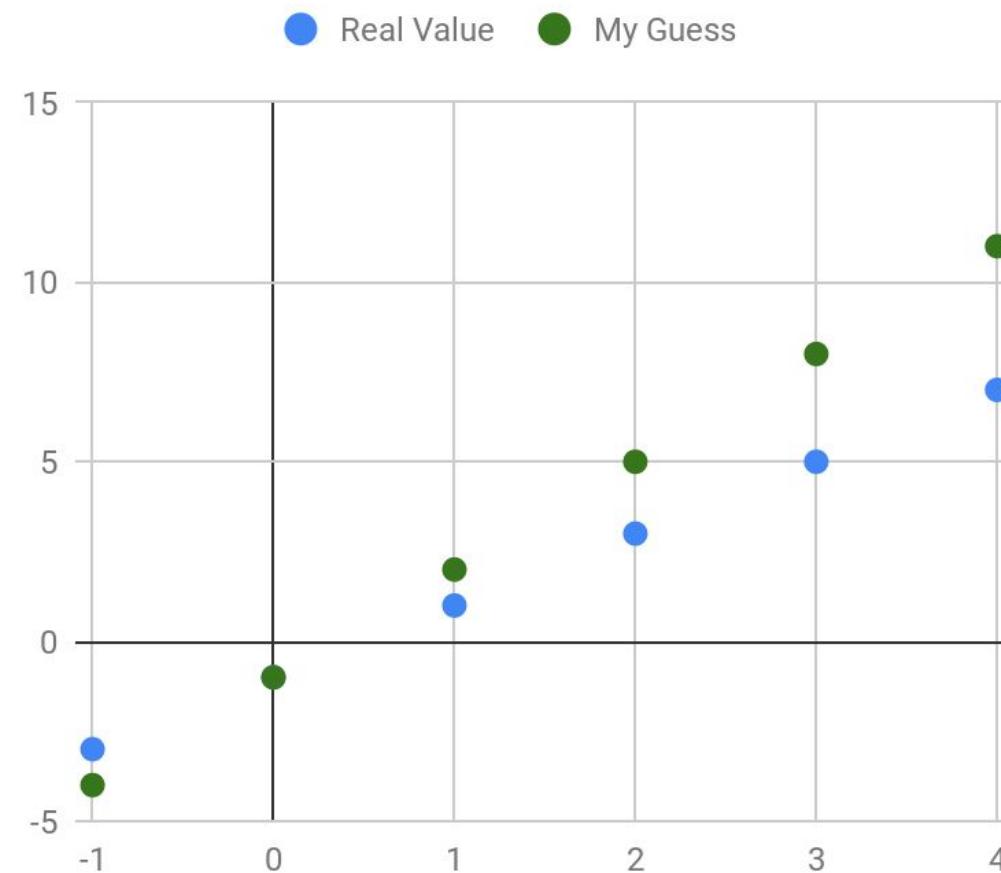
$$Y = 3X - 1$$

$$X = \{ -1, 0, 1, 2, 3, 4 \}$$

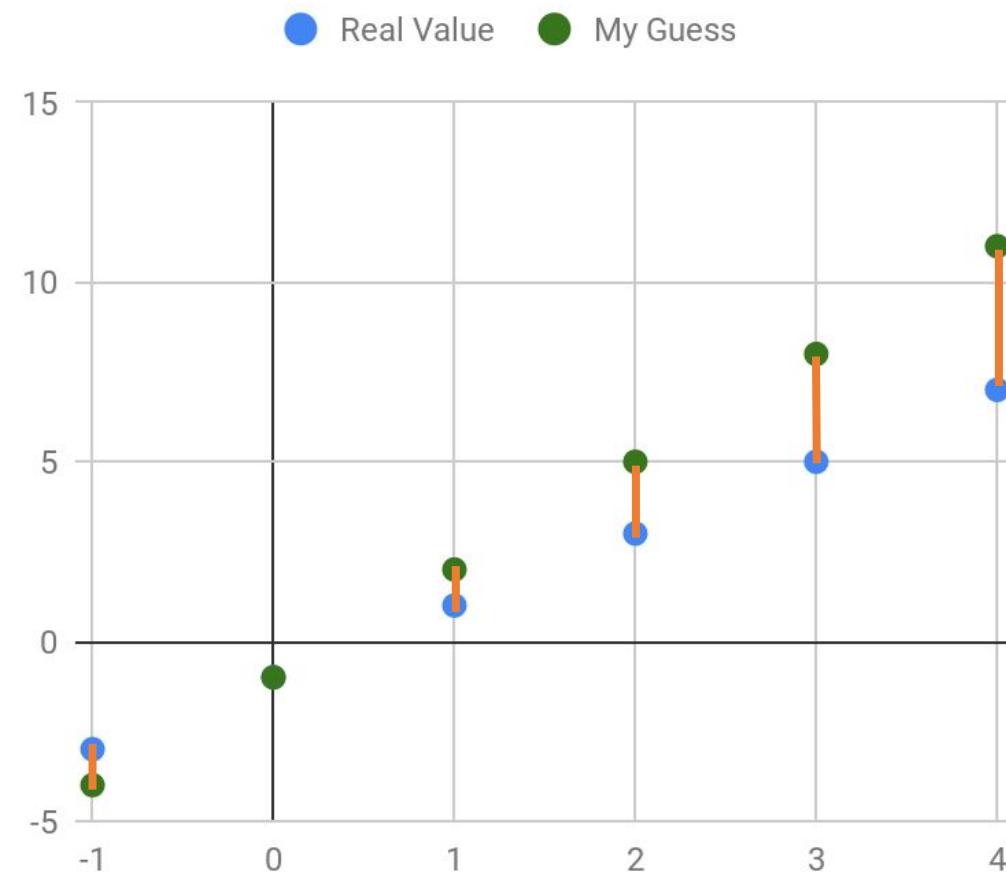
$$\text{My } Y = \{ -4, -1, 2, 5, 8, 11 \}$$

$$\text{Real } Y = \{ -3, -1, 1, 3, 5, 7 \}$$

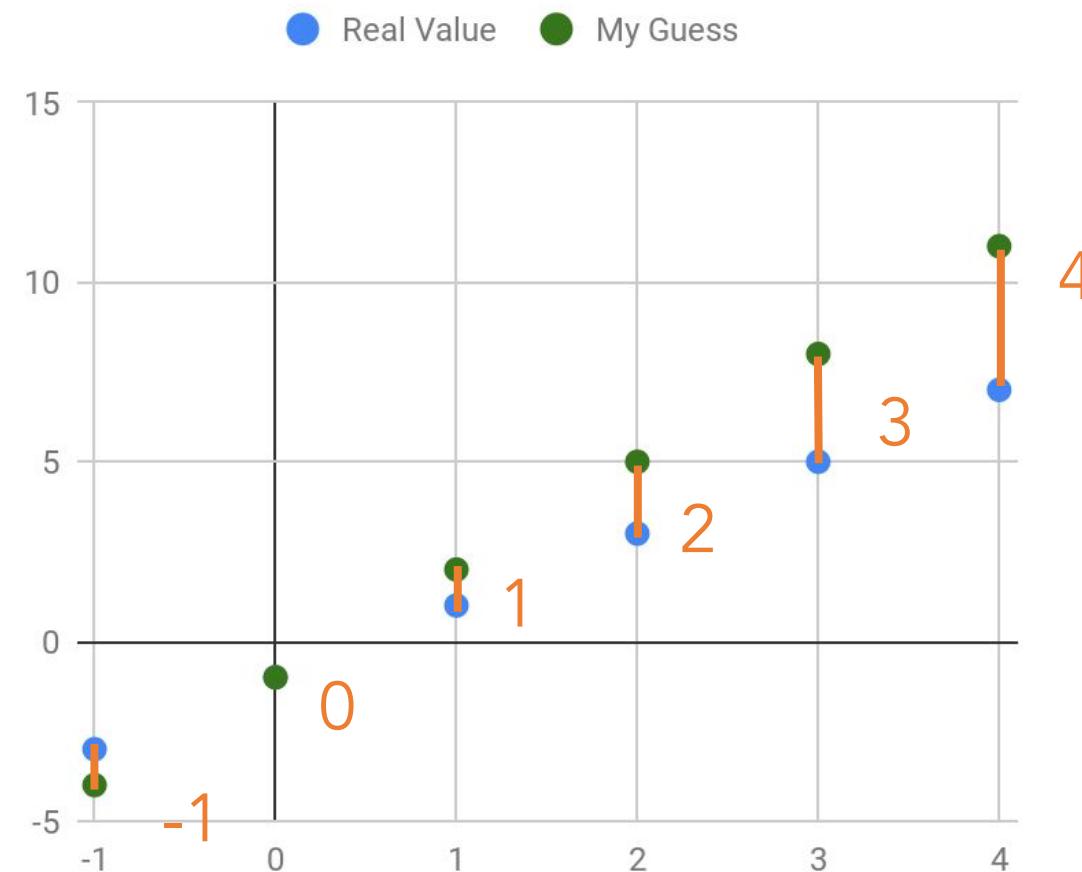
Let's measure it!



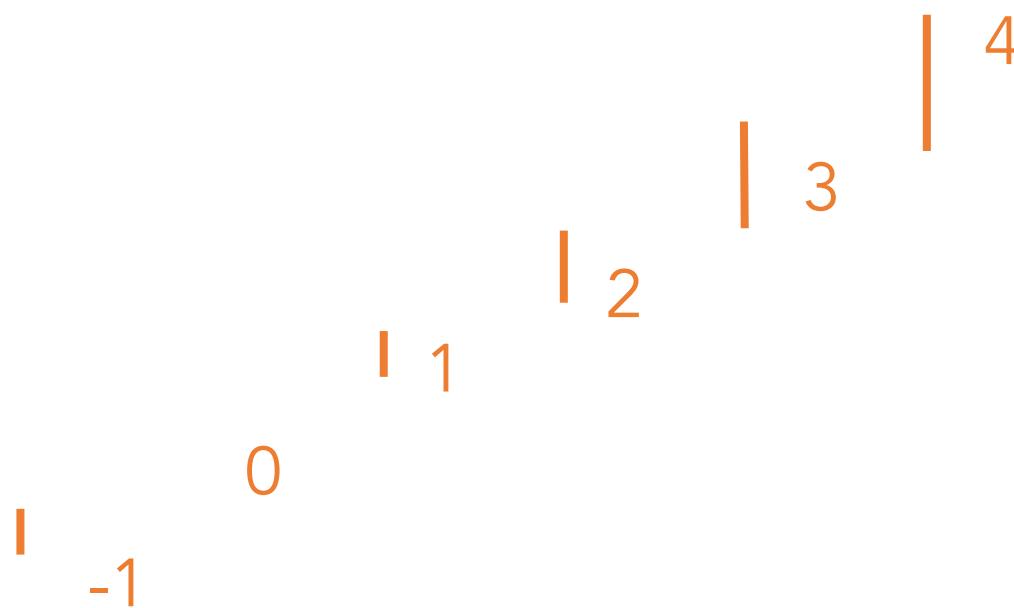
Let's measure it!



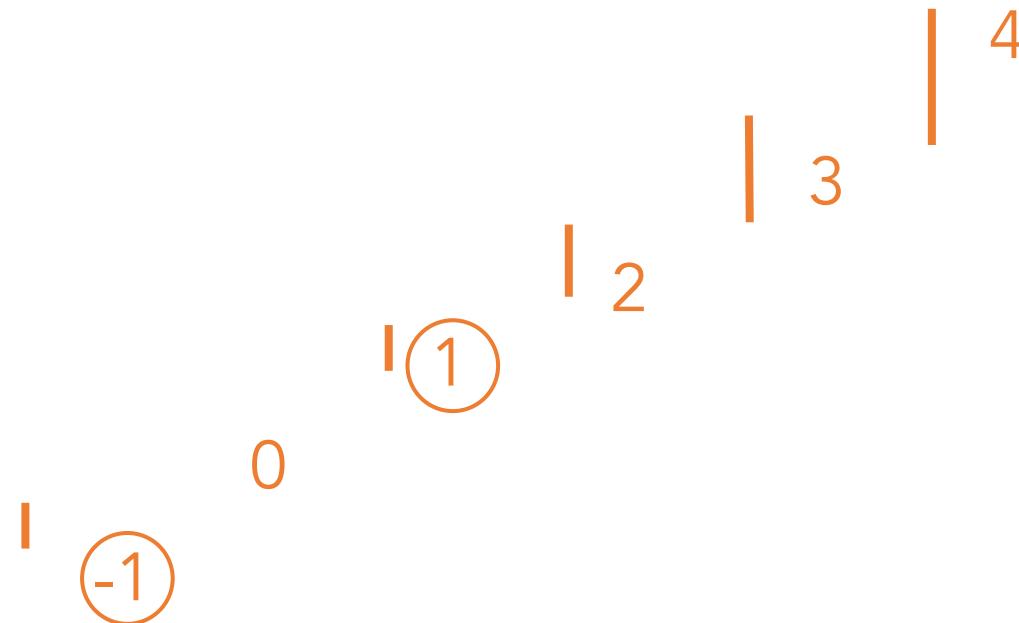
Let's measure it!



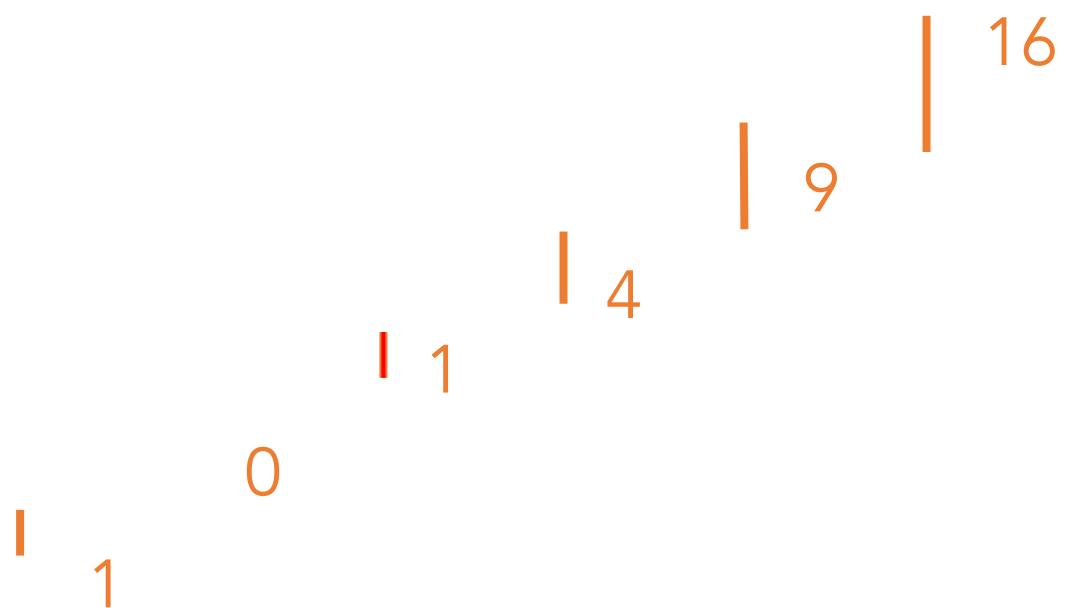
Let's measure it!



Houston, we have a
problem!



What if we square² them?

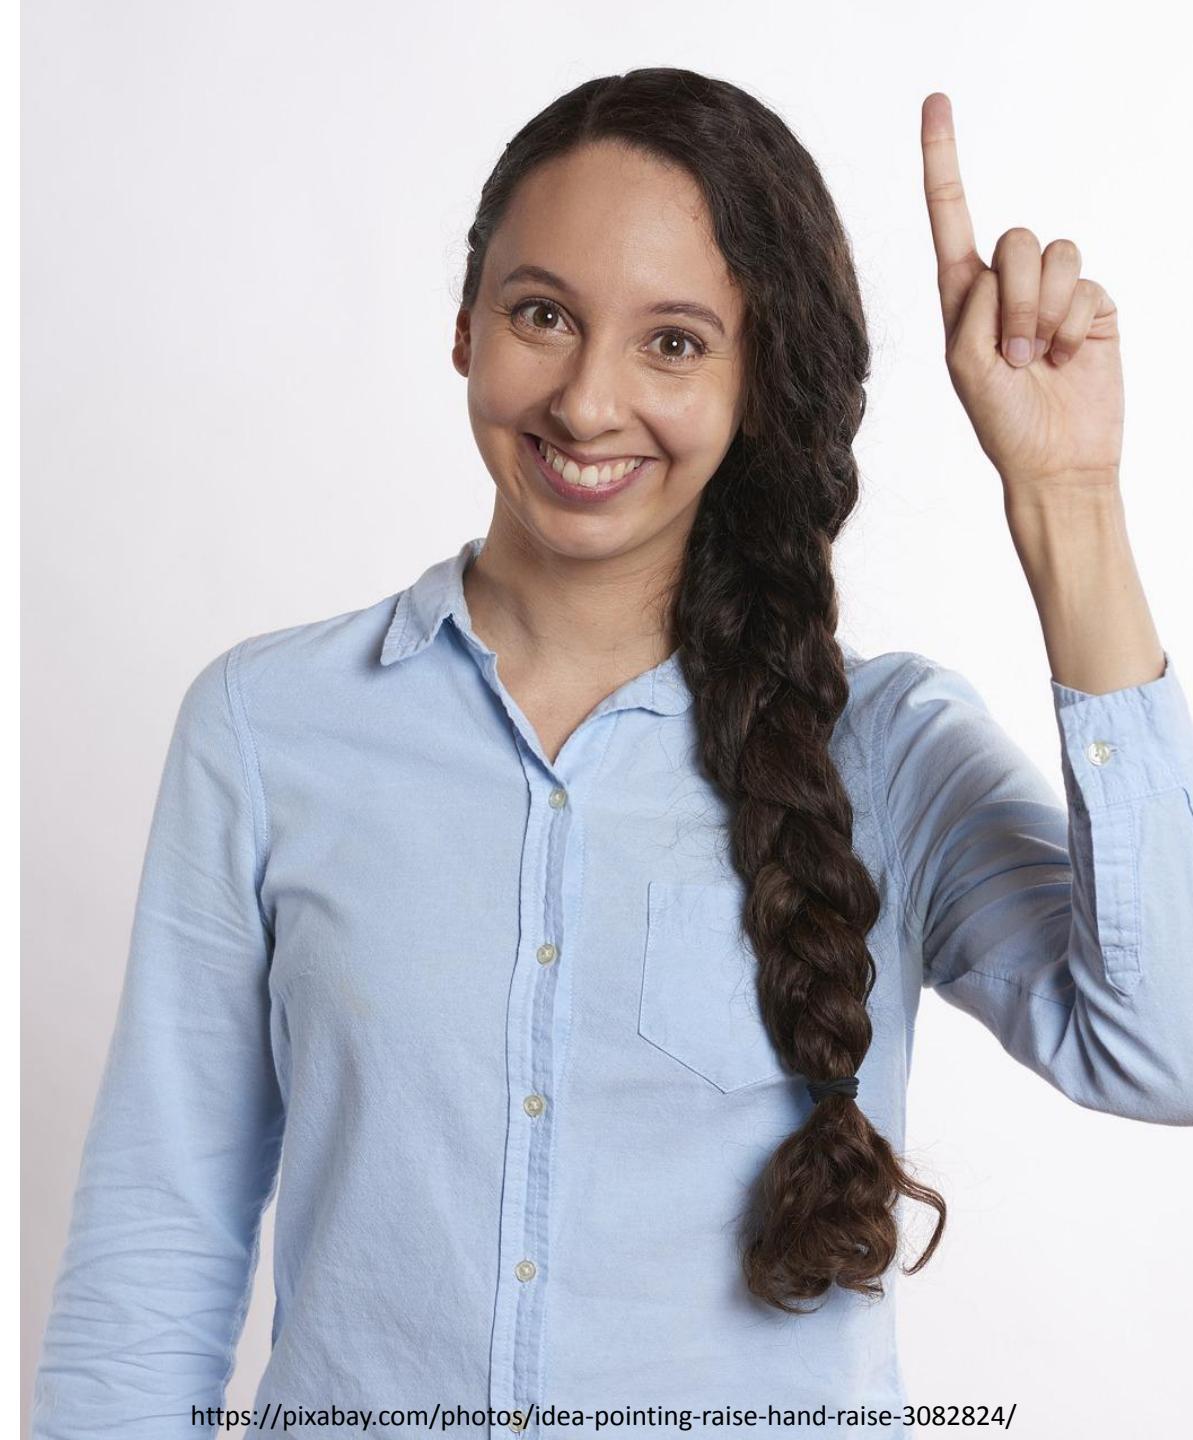


Total that (Σ) and take
the square root $\sqrt{\quad}$

$$\sqrt{1 + 1 + 4 + 9 + 16}$$

$$= \sqrt{31}$$

$$= 5.57$$



Make another guess!

$$Y = 2X - 2$$

$$X = \{ -1, 0, 1, 2, 3, 4 \}$$

$$\text{My } Y = \{ -4, -2, 0, 2, 4, 6 \}$$

$$\text{Real } Y = \{ -3, -1, 1, 3, 5, 7 \}$$

$$\text{Diff}^2 = \{ 1, 1, 1, 1, 1 \}$$

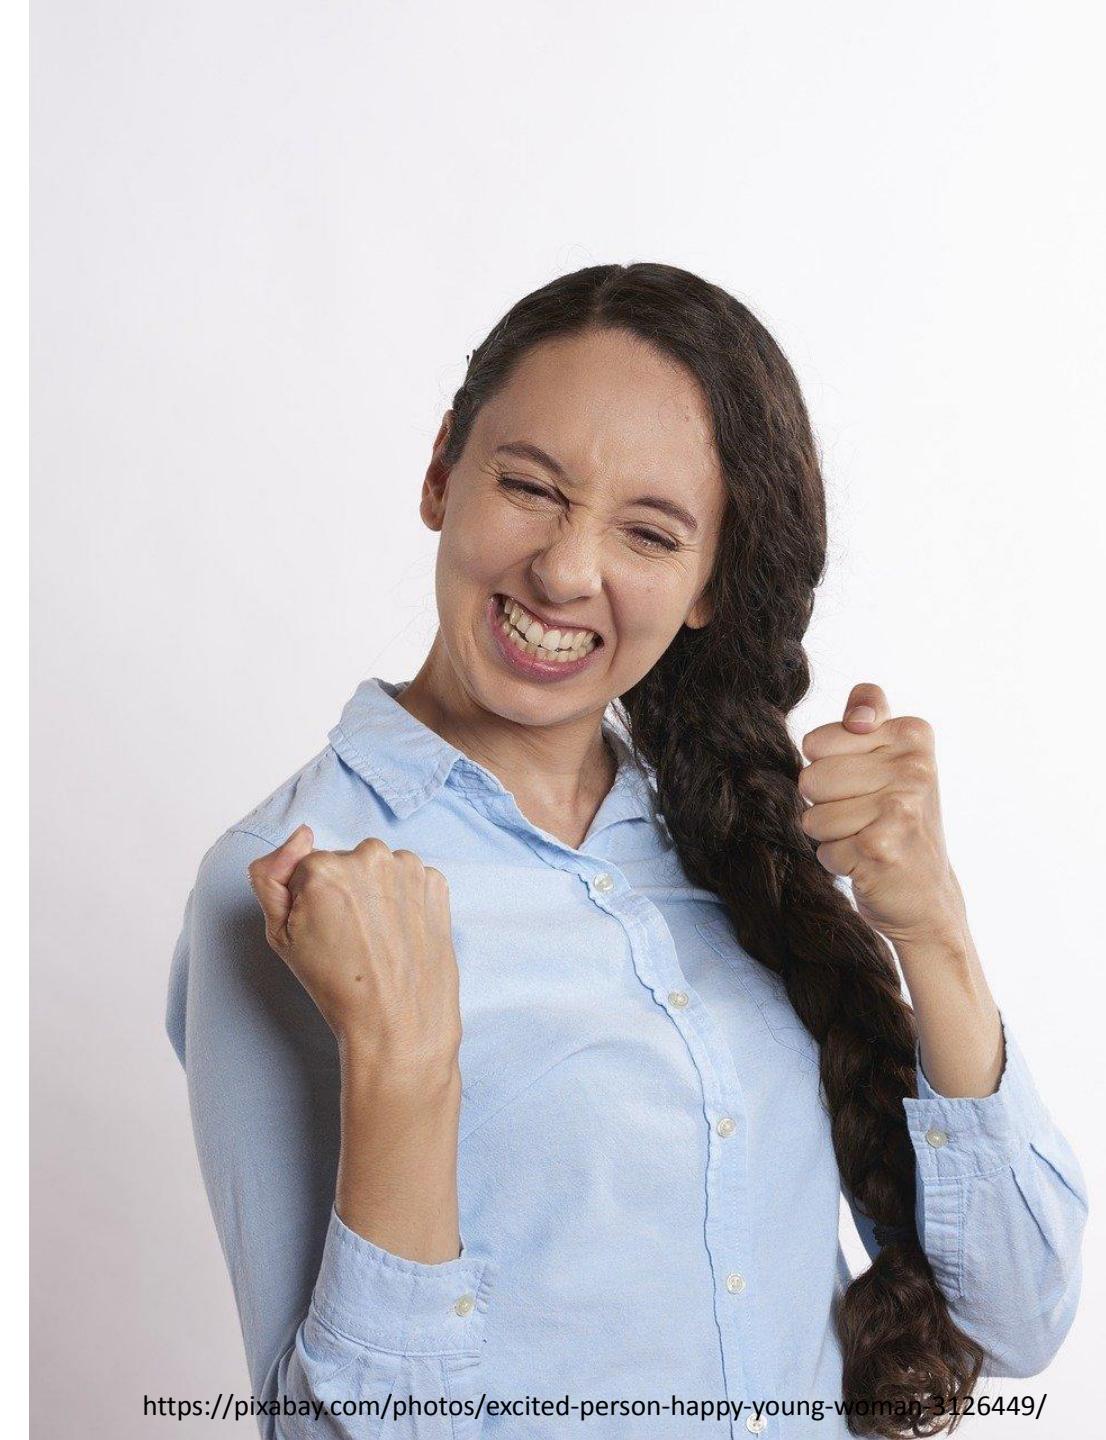


Get the same difference, repeat the same process.

$$\sqrt{1 + 1 + 1 + 1 + 1}$$

$$= \sqrt{5}$$

$$= 2.23$$



Make another guess!

$$Y = 2X - 1$$

$$X = \{ -1, 0, 1, 2, 3, 4 \}$$

$$\text{My } Y = \{ -3, -1, 1, 3, 5, 7 \}$$

$$\text{Real } Y = \{ -3, -1, 1, 3, 5, 7 \}$$

$$\text{Diff}^2 = \{ 0, 0, 0, 0, 0 \}$$



Make another guess!

$$Y = 2X - 1$$

$$X = \{-1, 0, 1, 2, 3, 4\}$$

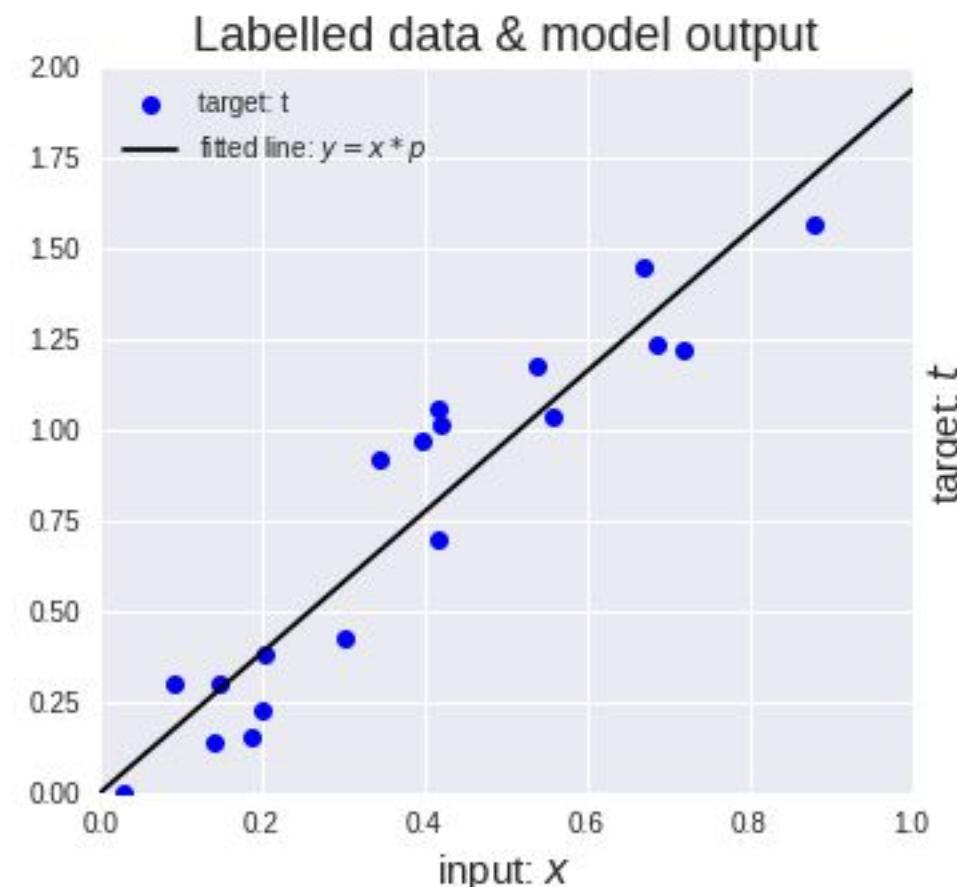
$$\text{My } Y = \{-3, -1, 1, 3, 5, 7\}$$

$$\text{Real } Y = \{-3, -1, 1, 3, 5, 7\}$$

$$\text{Diff}^2 = \{0, 0, 0, 0, 0\}$$

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

Goal is to Minimize MSE (Mean Squared Error)



Finding out the best solution

Trial and error approach

Loss
Function

← Parameter →

Loss
Function

Minimum of
Loss Function



Loss
Function



Loss
Function

Gradient of
value



Loss
Function



Move in Direction of Gradient
Learning Rate is size of the step to take



Loss
Function

End up here



Loss
Function

Get the
gradient



Loss
Function

Move in Direction of Gradient



Loss
Function

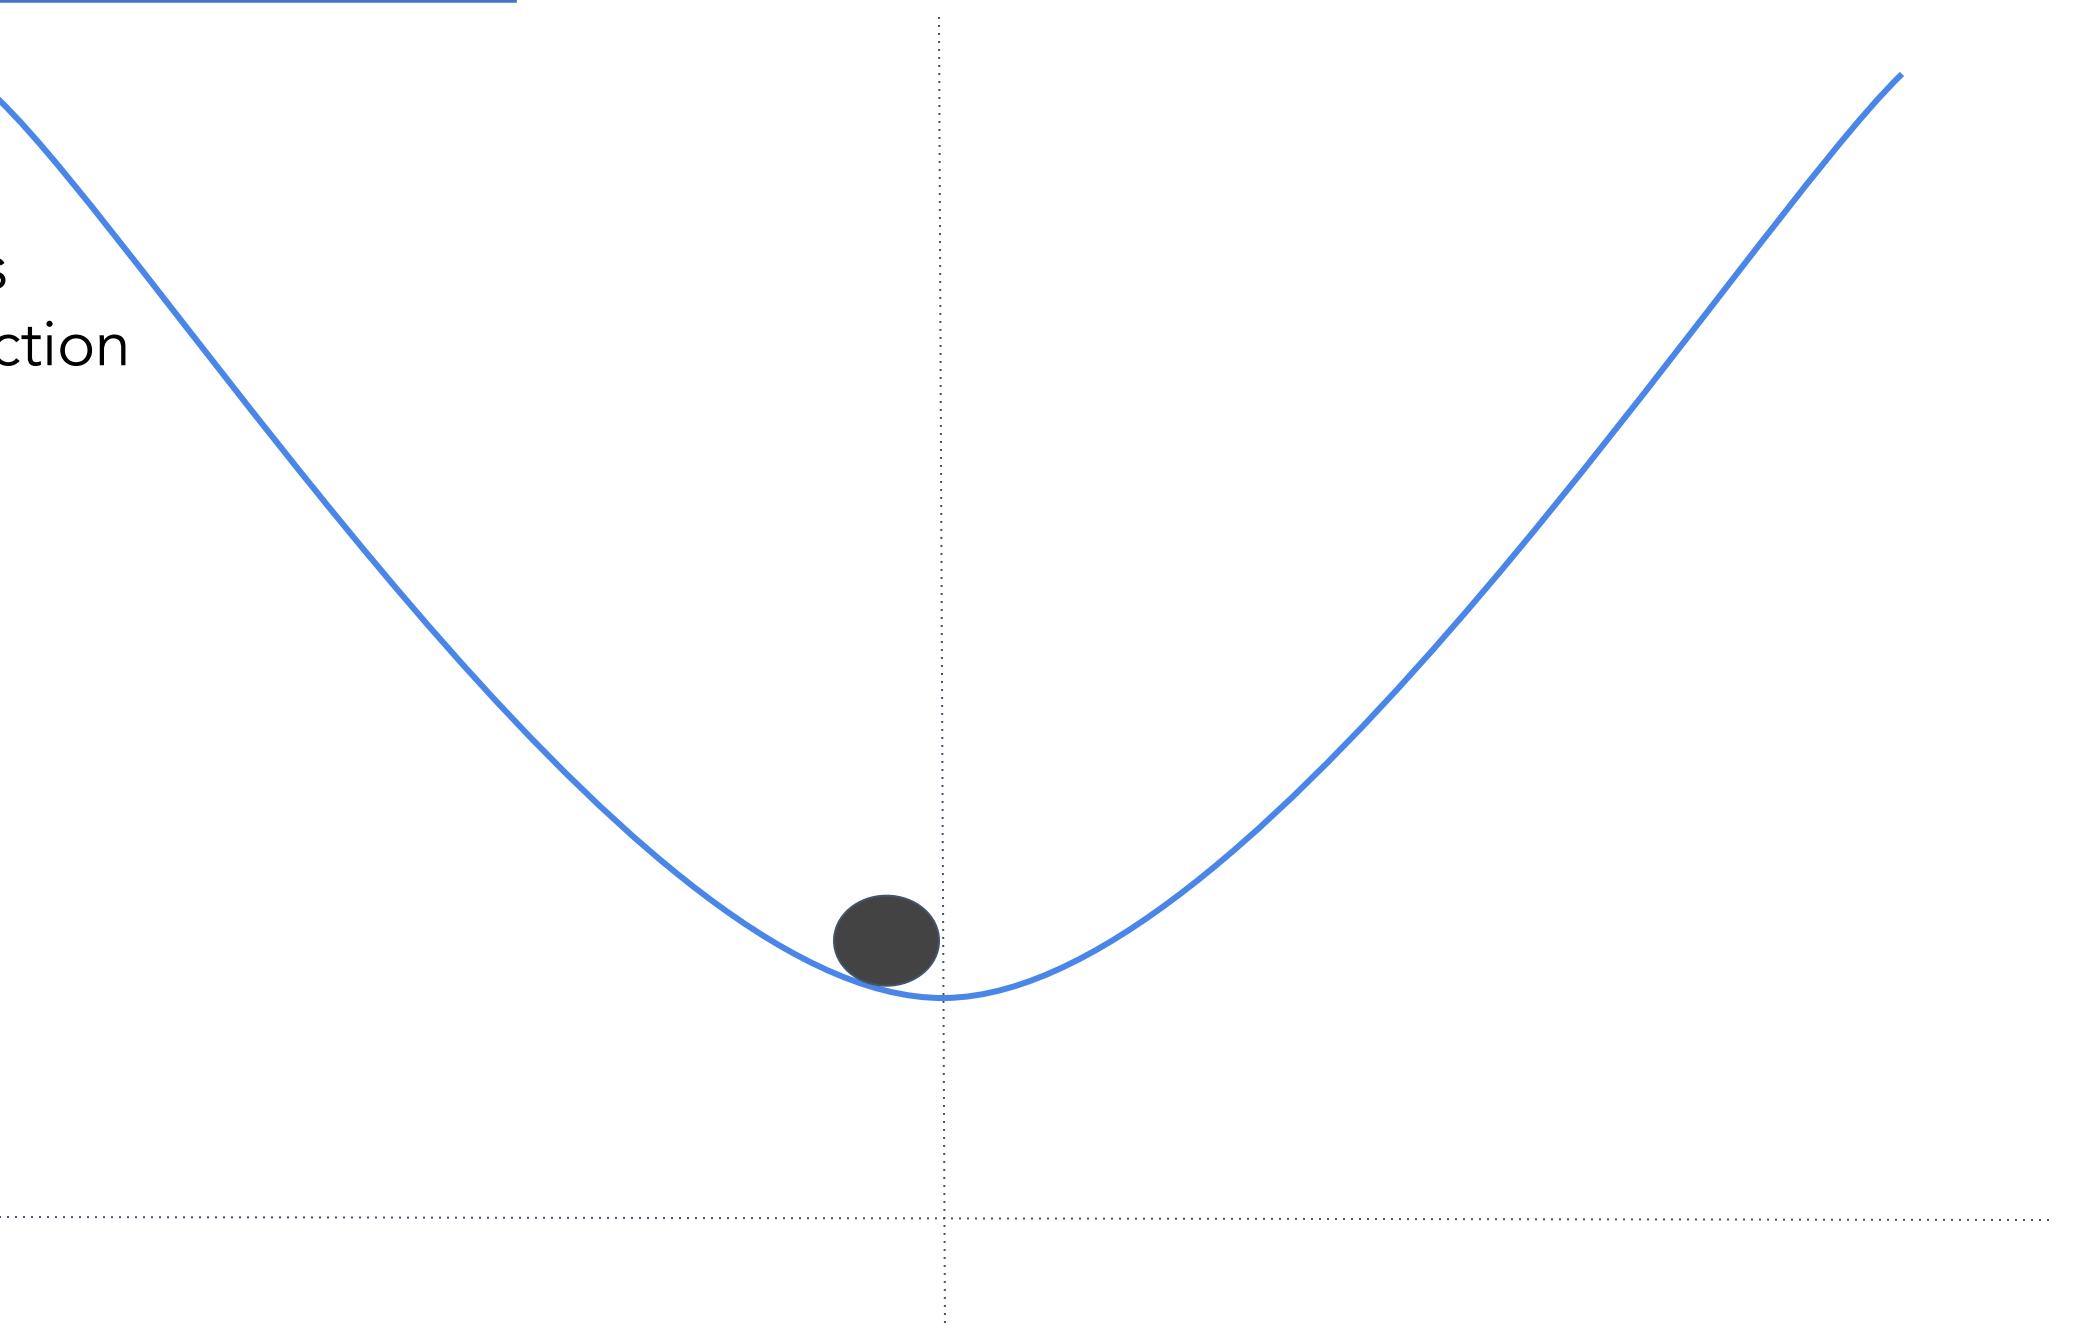
End Up here



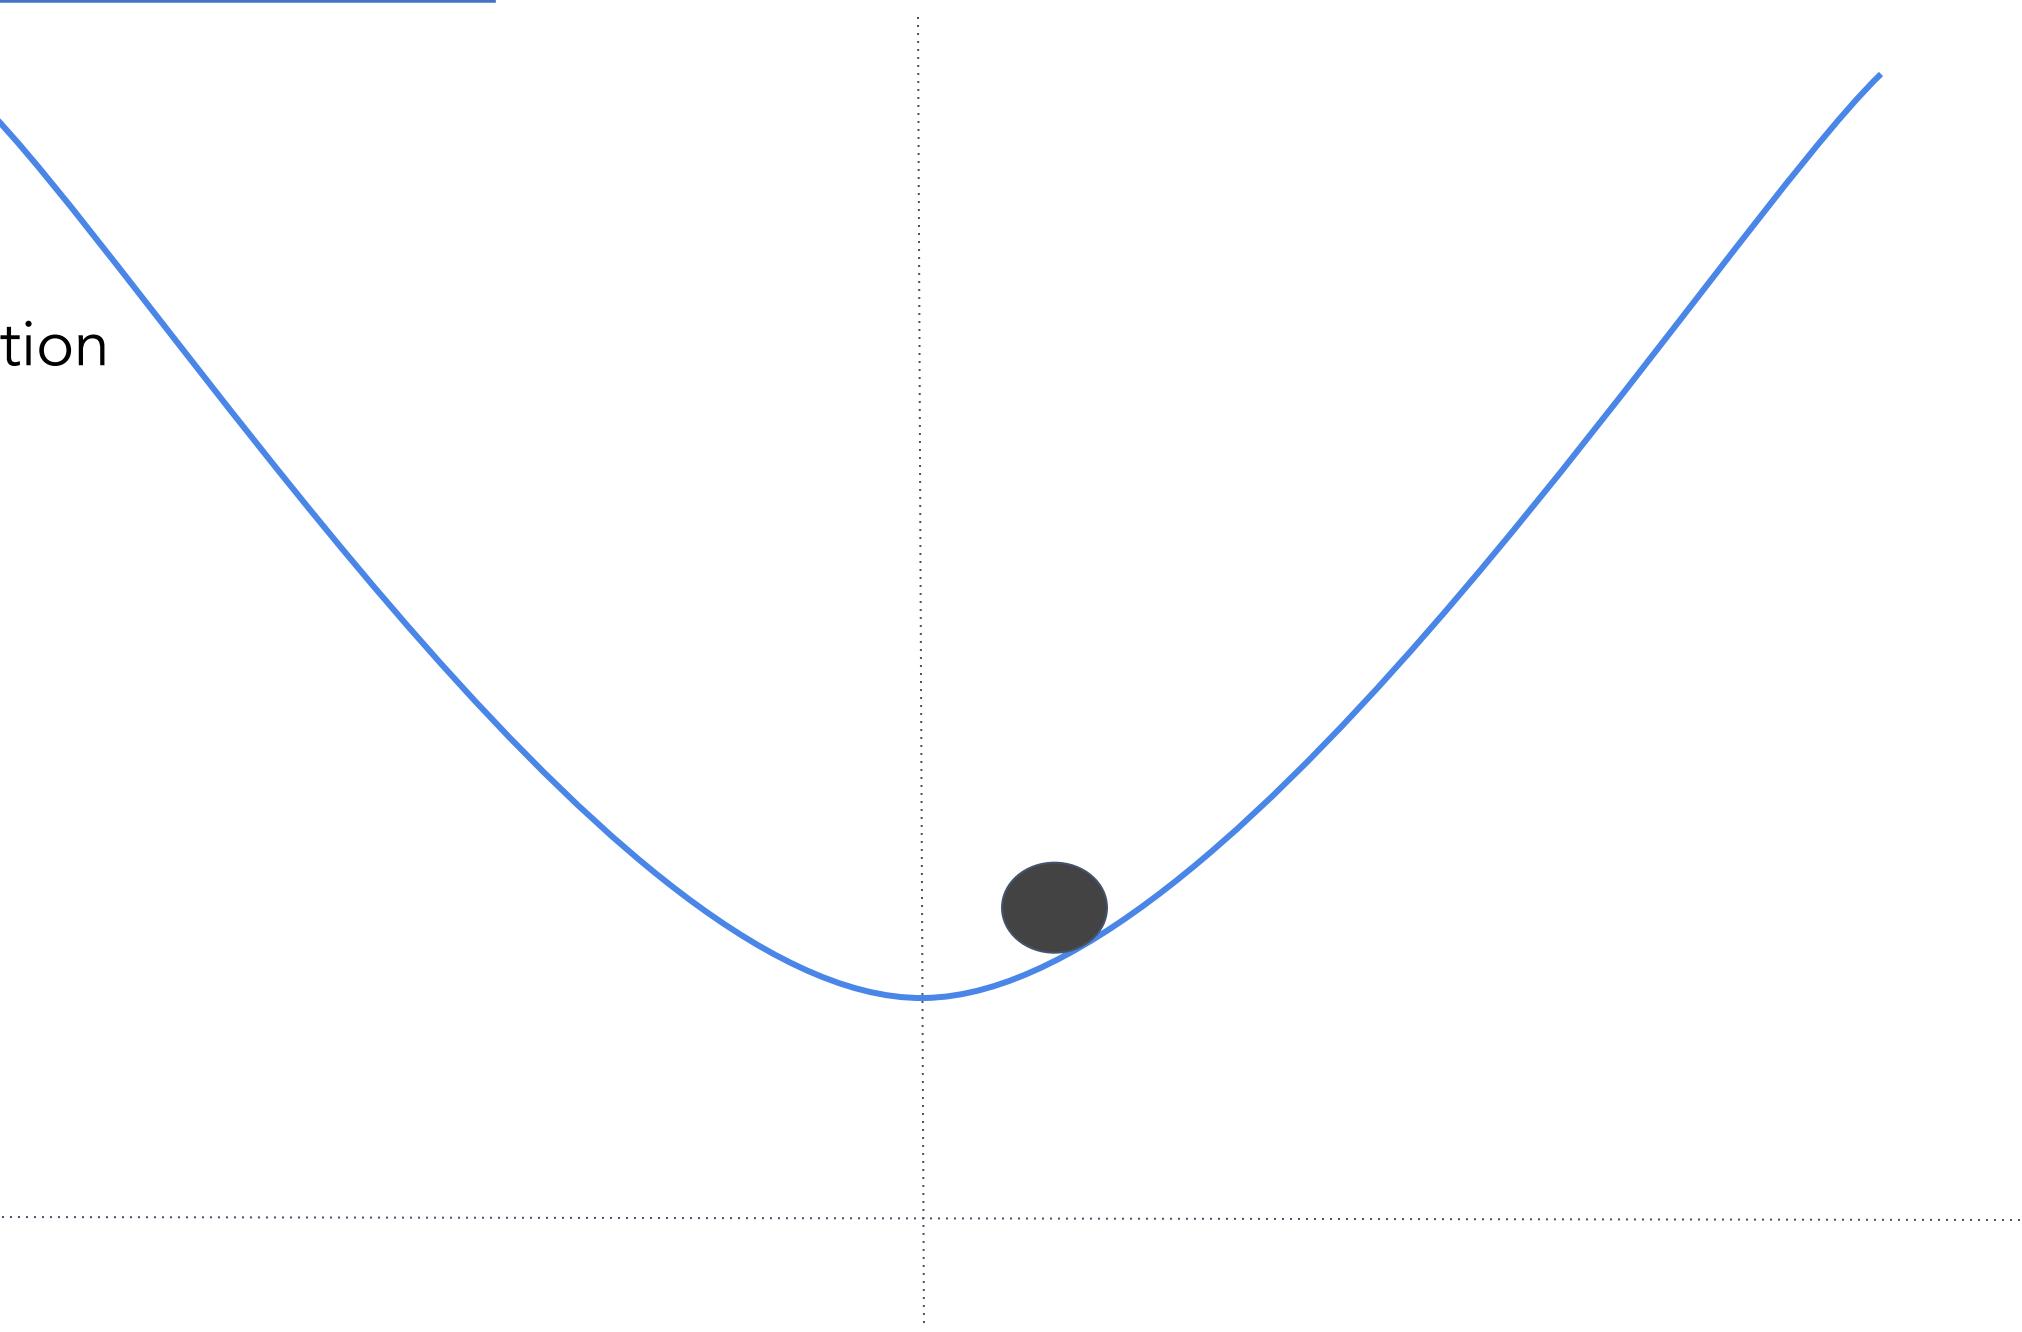
Loss
Function



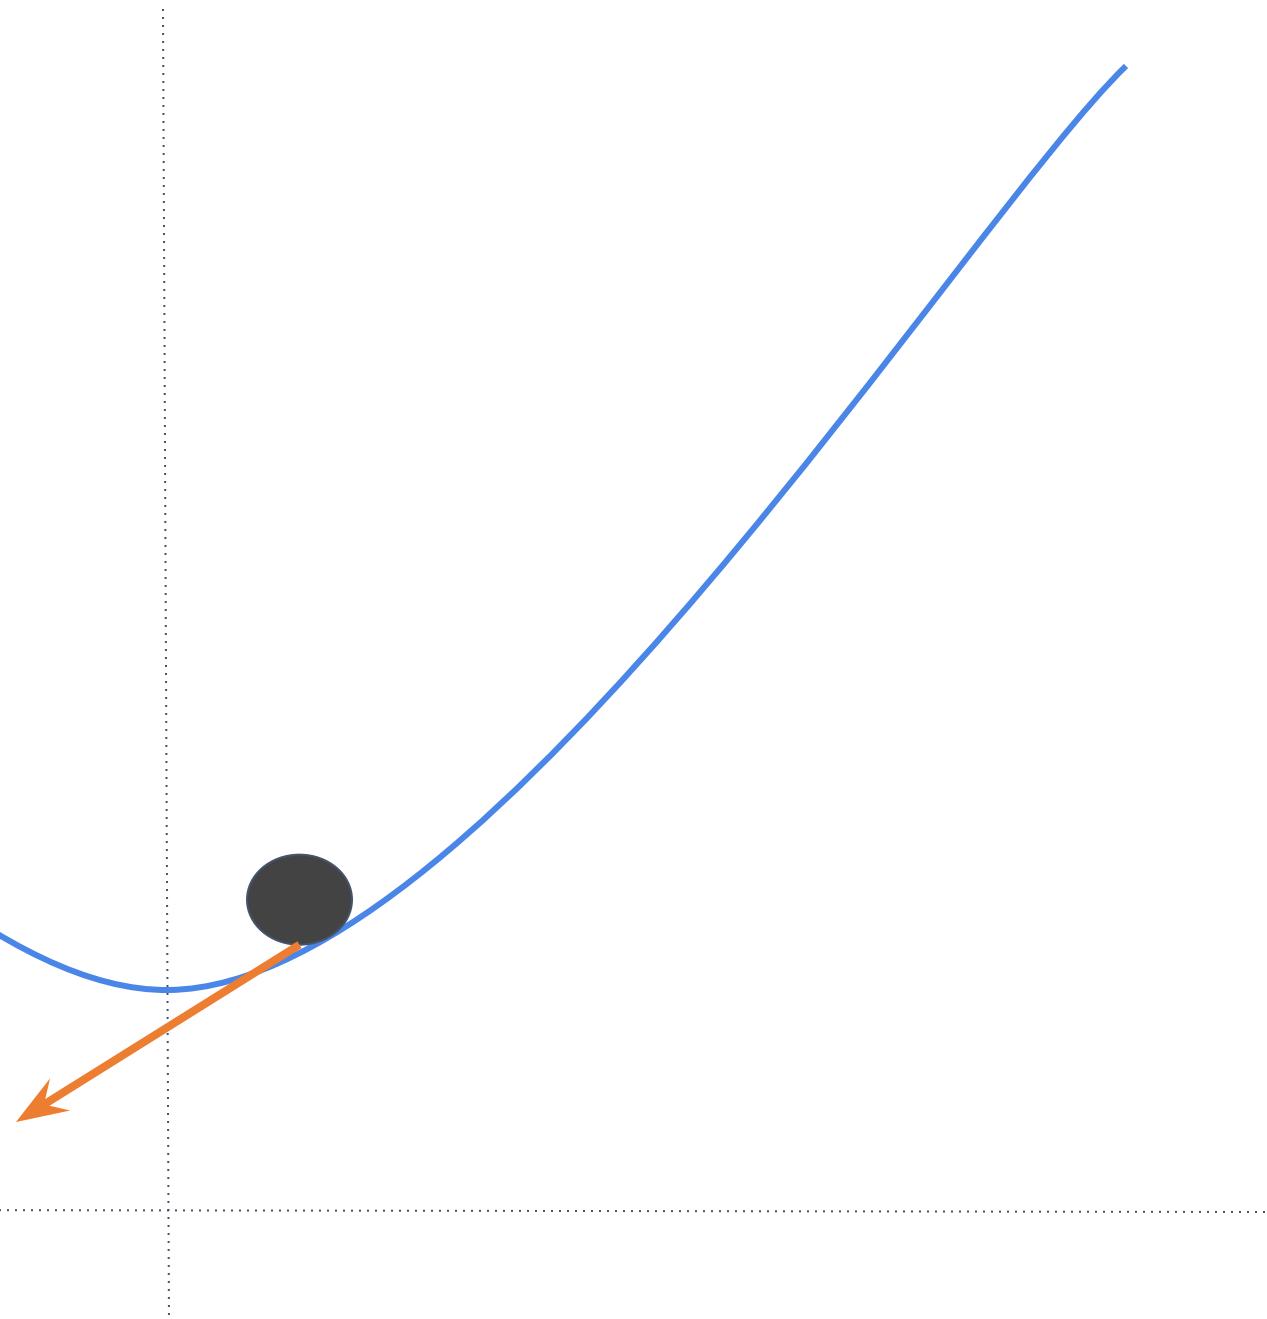
Loss
Function



Loss
Function



Loss
Function



Loss
Function



Loss
Function



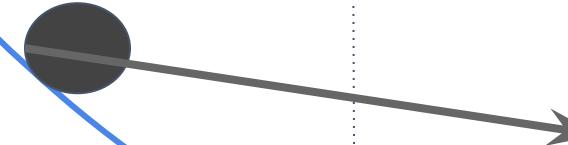
Loss
Function



Loss
Function



Loss
Function



Loss
Function

Move in Direction of Gradient



Loss
Function

Move in Direction of Gradient

Loss
Function

Move in Direction of Gradient

Loss
Function

Move in Direction of Gradient



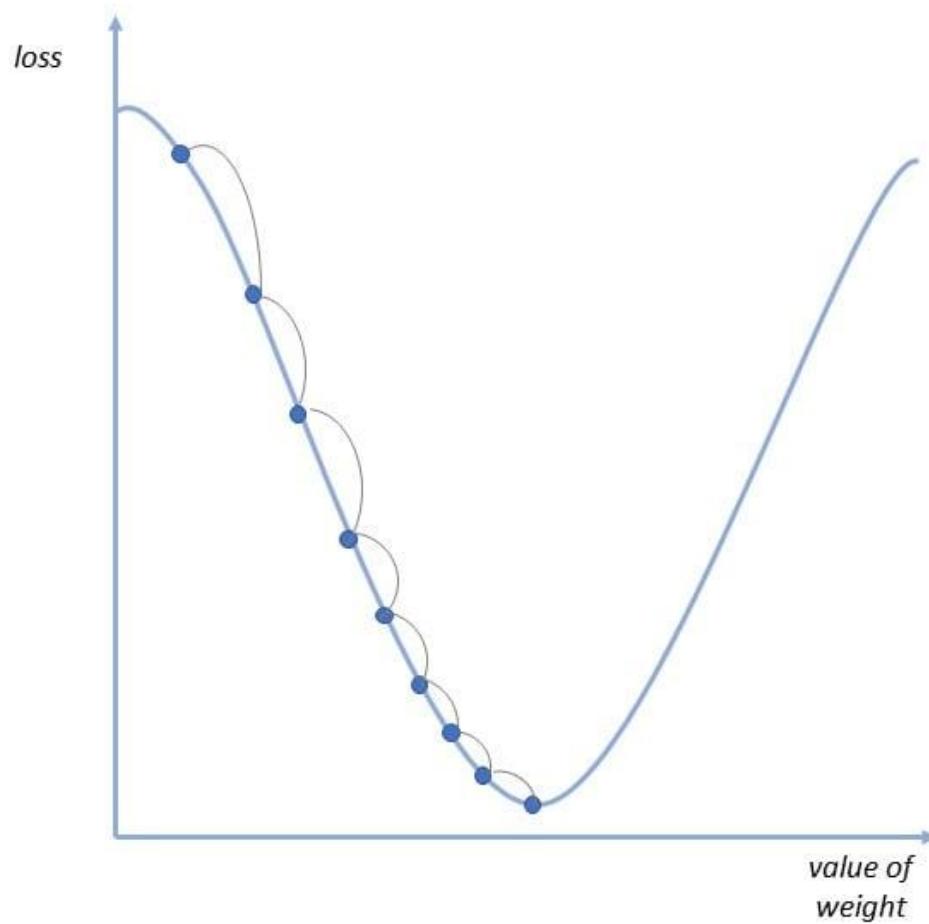
Loss
Function

Move in Direction of Gradient

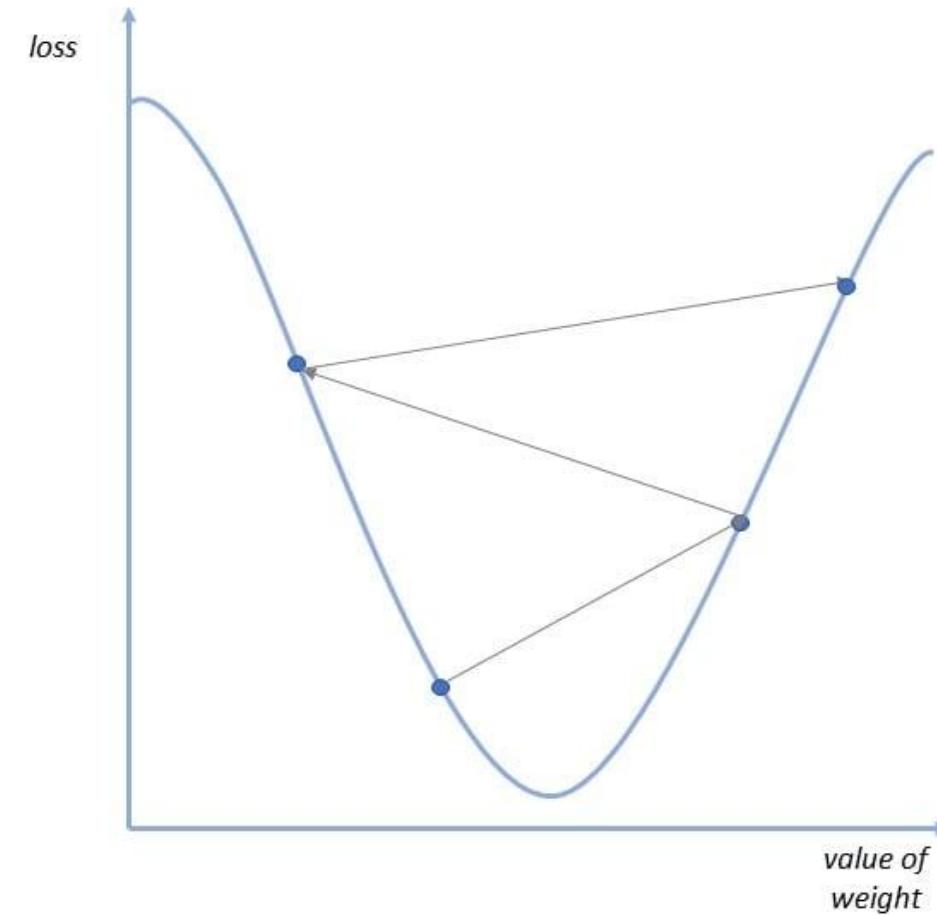


It is important to choose the correct Learning Rate (size of the step)

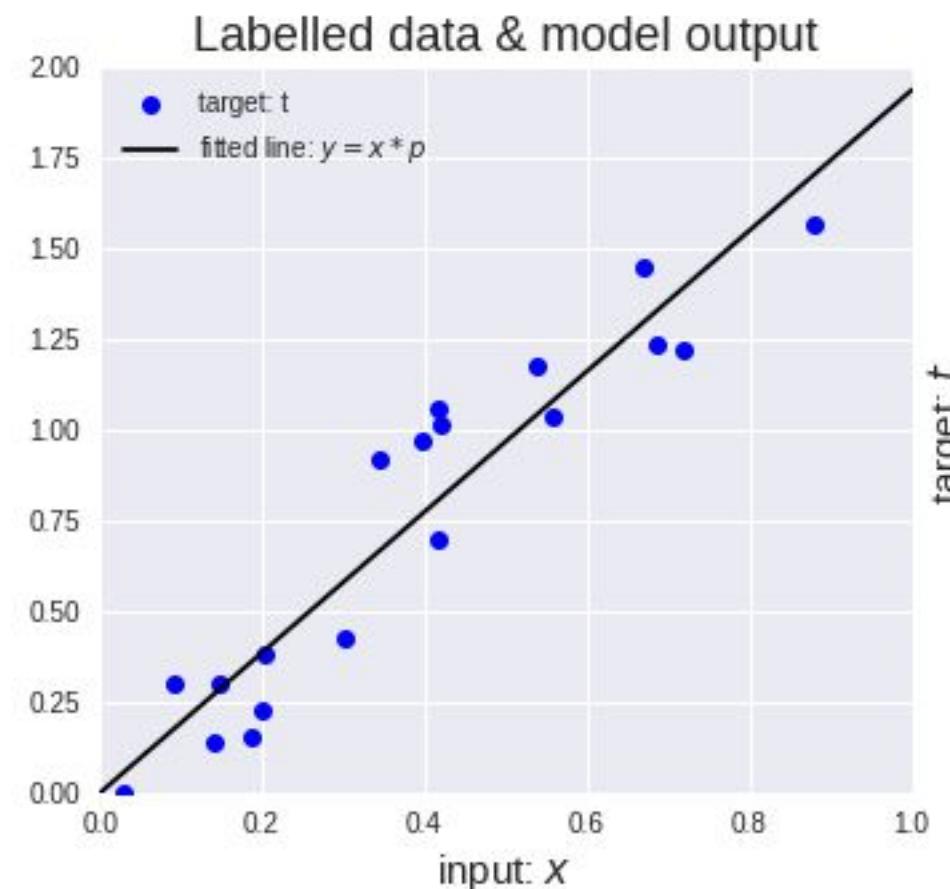
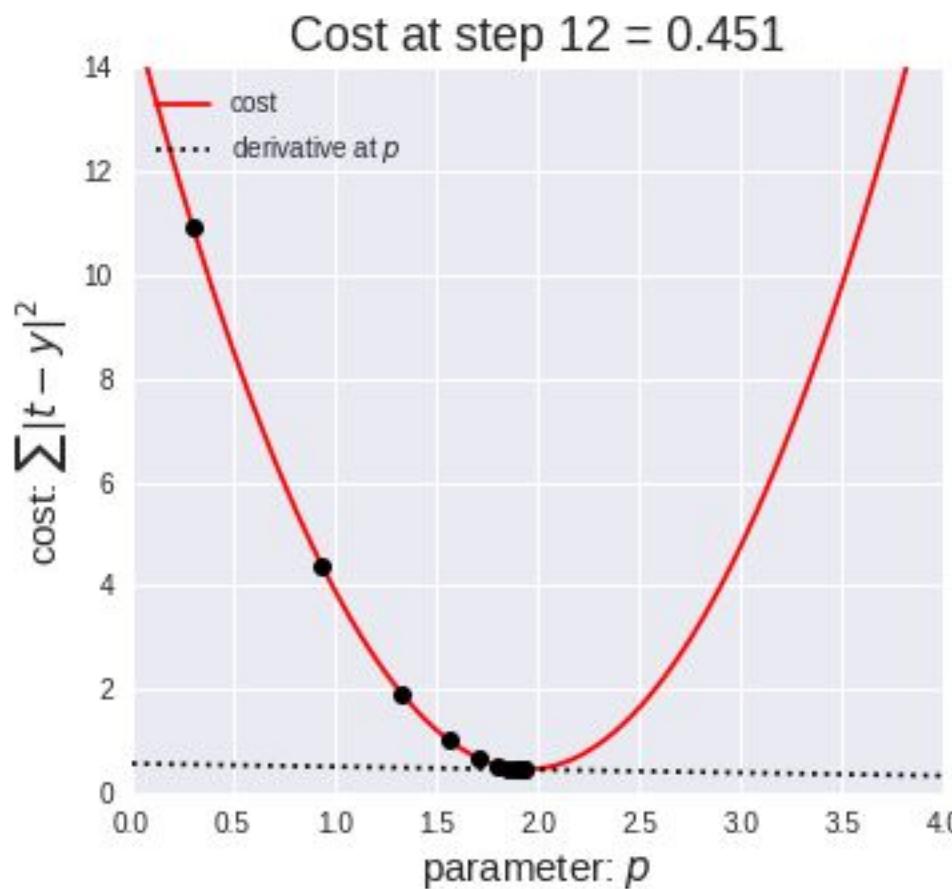
If the **Learning Rate** is too small it may take a long time to reach the minimum



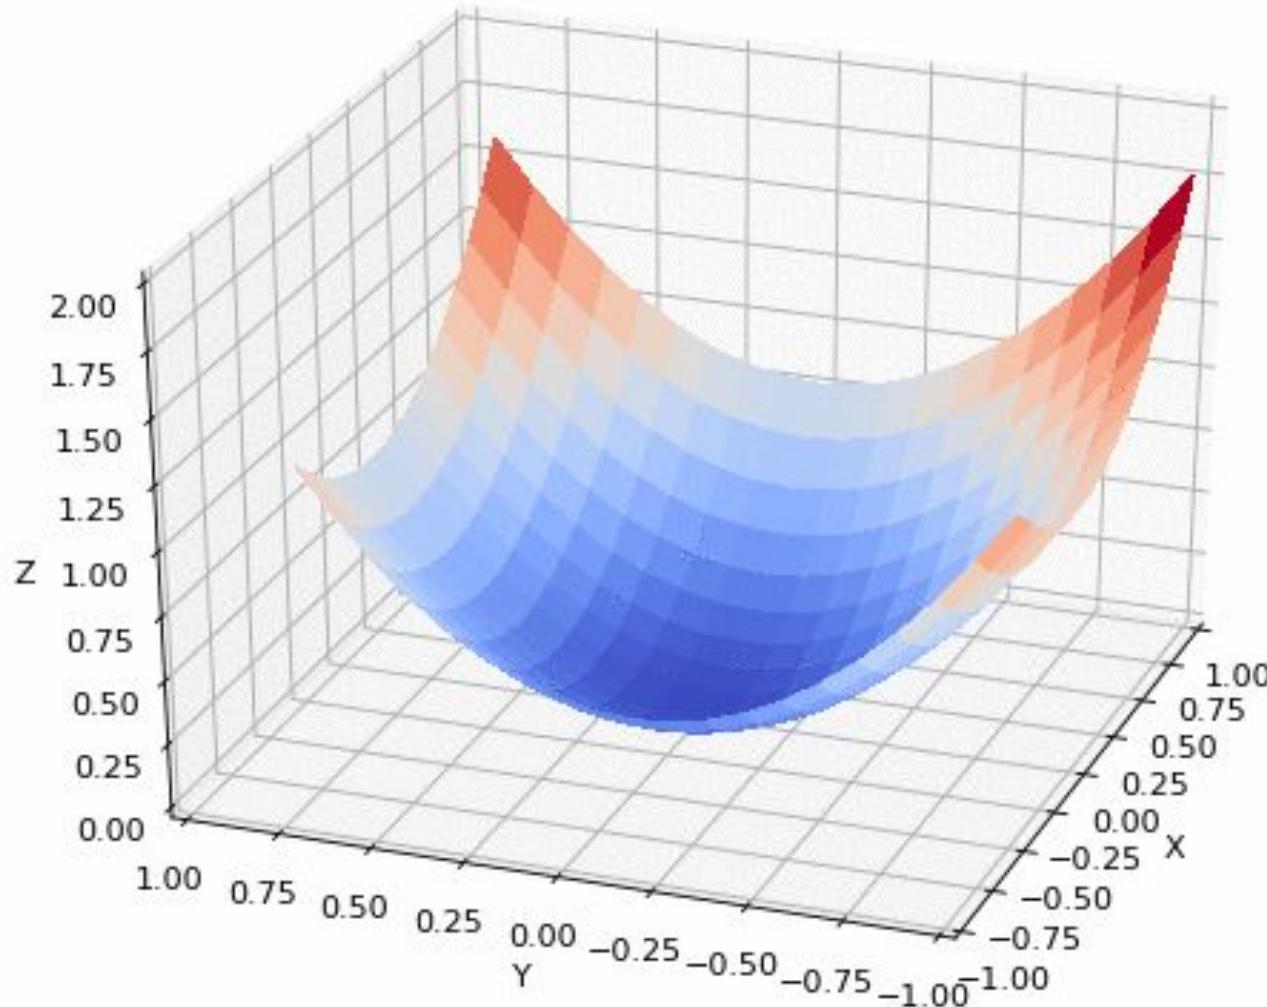
If the **Learning Rate** is too large we may never reach the minimum



Gradient Descent algorithm

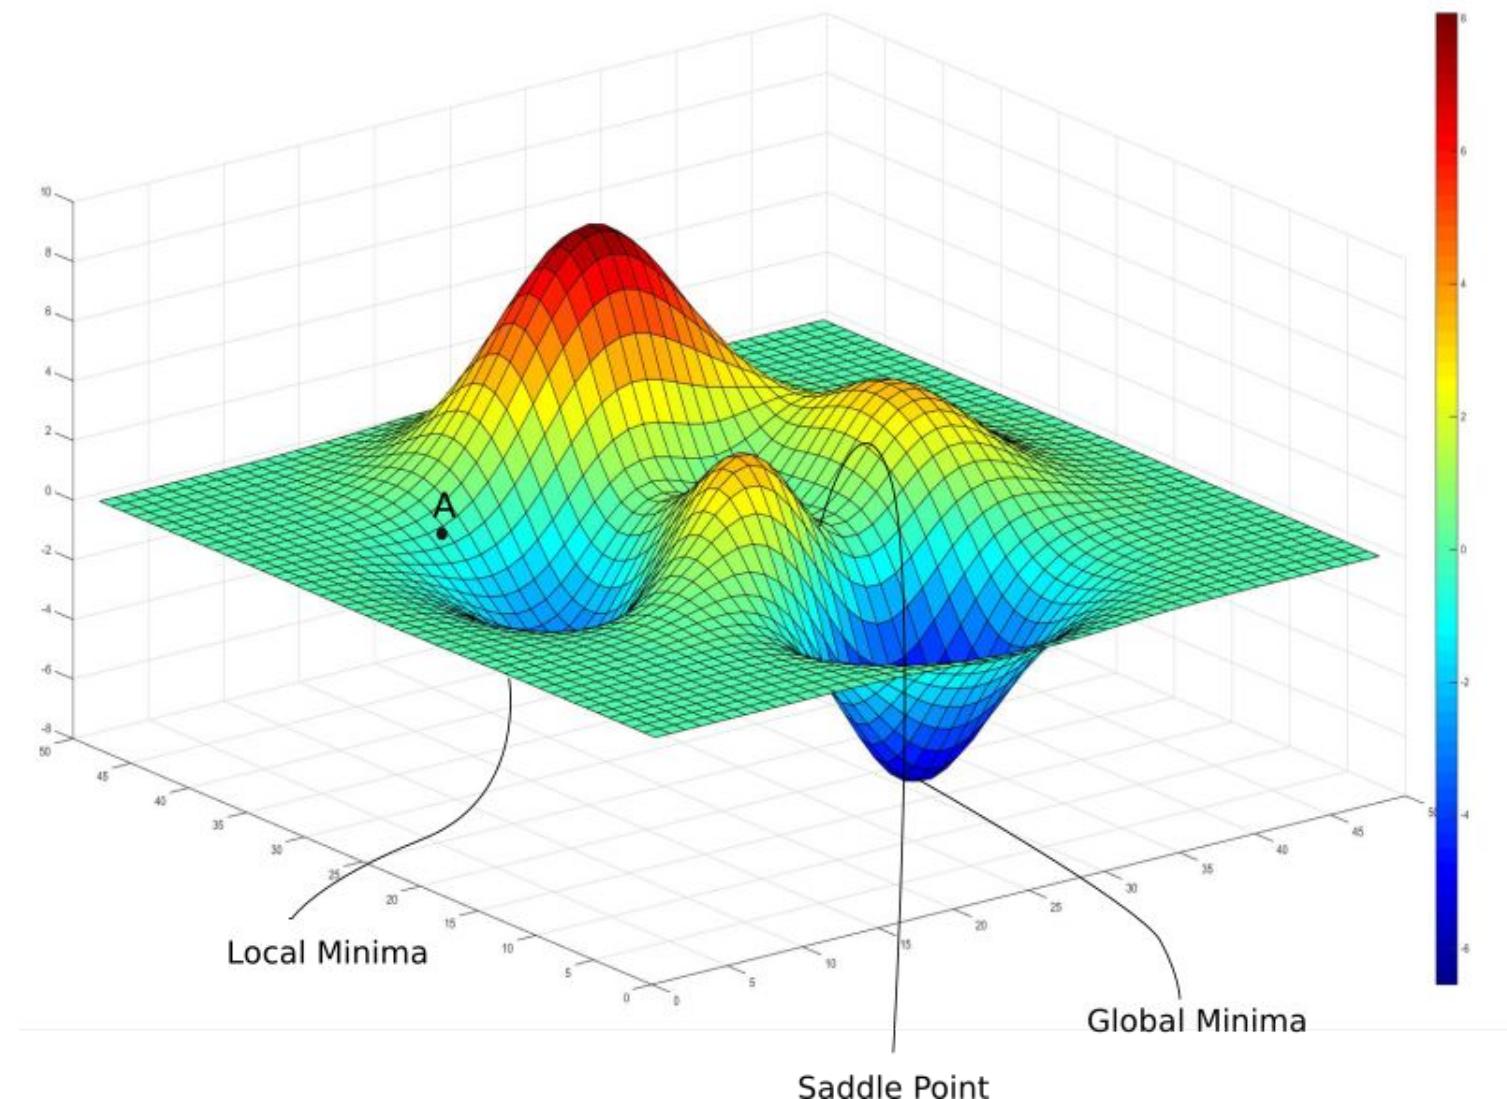


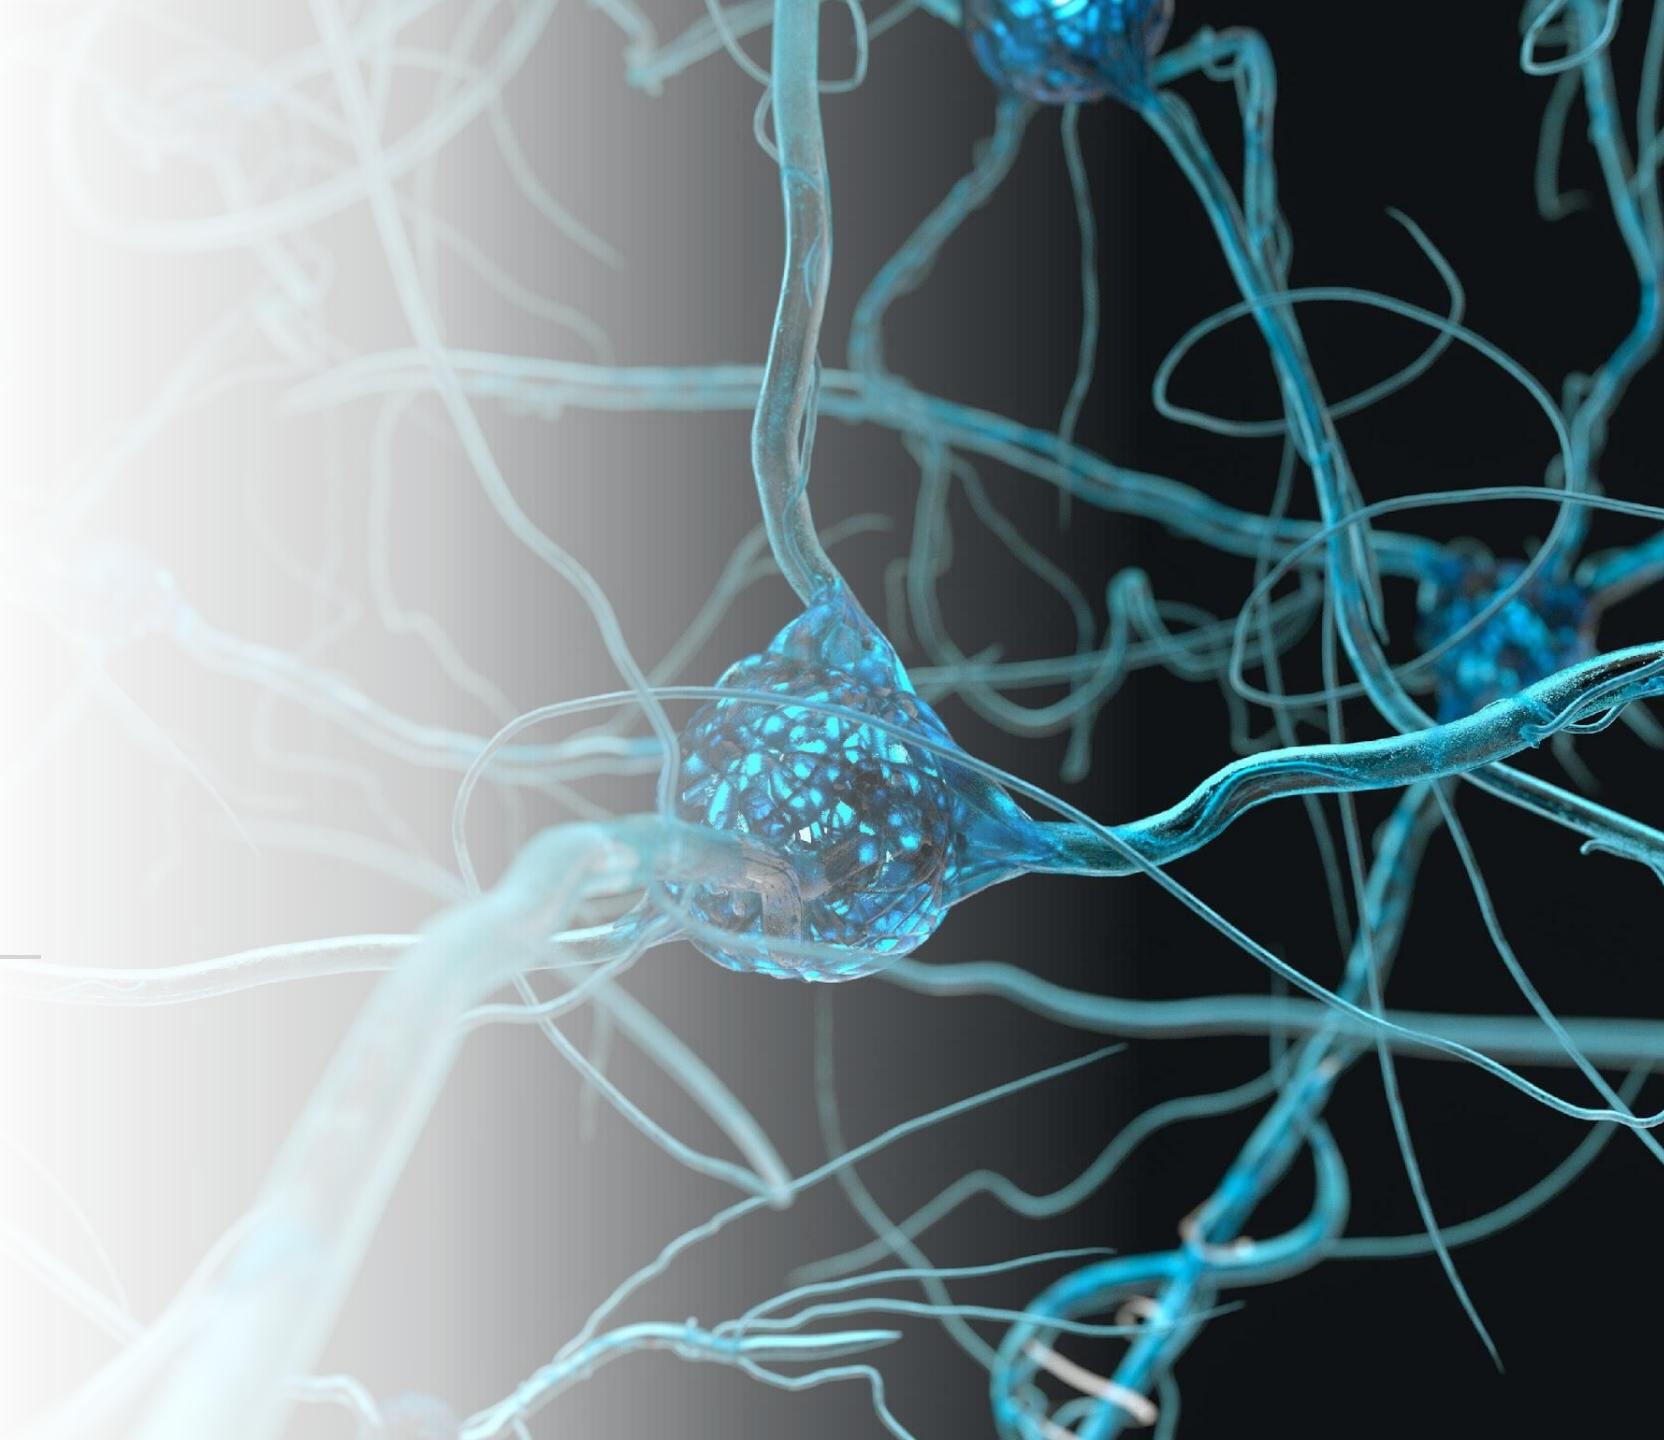
Gradient Descent for Two Parameters



A single minima
Global minima

Gradient Descent for Two Parameters





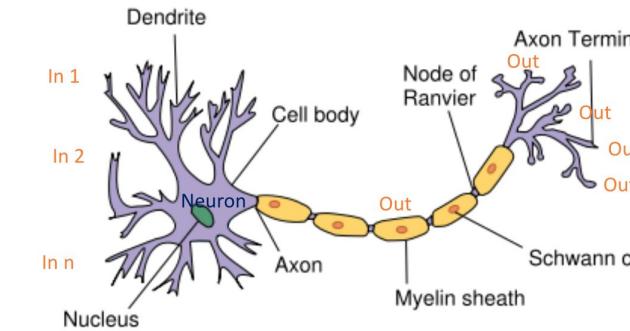
A horizontal orange bar is located at the top left of the slide.

Artificial Neural Networks

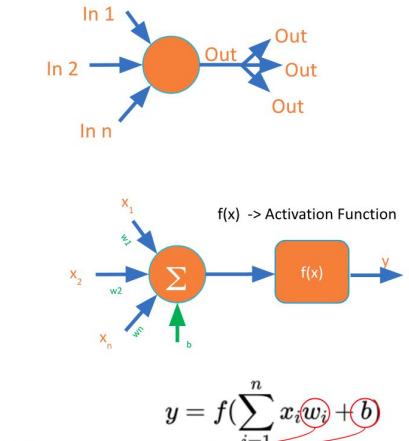
What is an Artificial Neural Network (ANN)?

Sequential

Neuron (Perceptron)

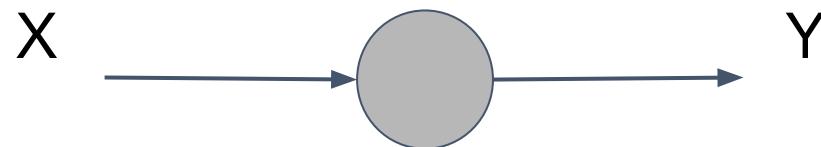


Dense Neural Network (DNN)



A neuron

a neuron's output is a function of its inputs (in this case only one)



$$y = f(x) = \mathbf{w}x + \mathbf{b}$$

There are only **two parameters** to adjust:
The **weight** for each input and a **bias**

First scenario: a regression

Linear Regression with a Single Neuron

colab.research.google.com
Regression.ipynb

```
✓ [2] import tensorflow as tf
      import numpy as np
      from tensorflow import keras
```

```
11s
# define a neural network with one neuron
# for more information on TF functions see: https://www.tensorflow.org/api\_docs
my_layer = keras.layers.Dense(units=1, input_shape=[1])
model = tf.keras.Sequential([my_layer])
```

```
# use stochastic gradient descent for optimization and
# the mean squared error loss function
model.compile(optimizer='sgd', loss='mean_squared_error')
```

```
# define some training data (xs as inputs and ys as outputs)
xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
```

```
# fit the model to the data (aka train the model)
model.fit(xs, ys, epochs=500)
```

1 layer, 1 neuron, 1 input

Stochastic gradient descent

Inputs and outputs (labels)

Train the model



Linear Regression with a Single Neuron

colab.research.google.com
Regression.ipynb

```
[2] import tensorflow as tf
import numpy as np
from tensorflow import keras

# define a neural network with one neuron
# for more information on TF functions see: https://www.tensorflow.org/api\_docs
my_layer = keras.layers.Dense(units=1, input_shape=[1])
model = tf.keras.Sequential([my_layer])

# use stochastic gradient descent for optimization and
# the mean squared error loss function
model.compile(optimizer='sgd', loss='mean_squared_error')

# define some training data (xs as inputs and ys as outputs)
xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)

# fit the model to the data (aka train the model)
model.fit(xs, ys, epochs=500)
```

```
Epoch 500/500
1/1 [=====] - 0s 6ms/step - loss: 3.4704e-05
<keras.callbacks.History at 0x7f1d6cccd7f10>
```

```
[4] print(model.predict([10.0]))
```

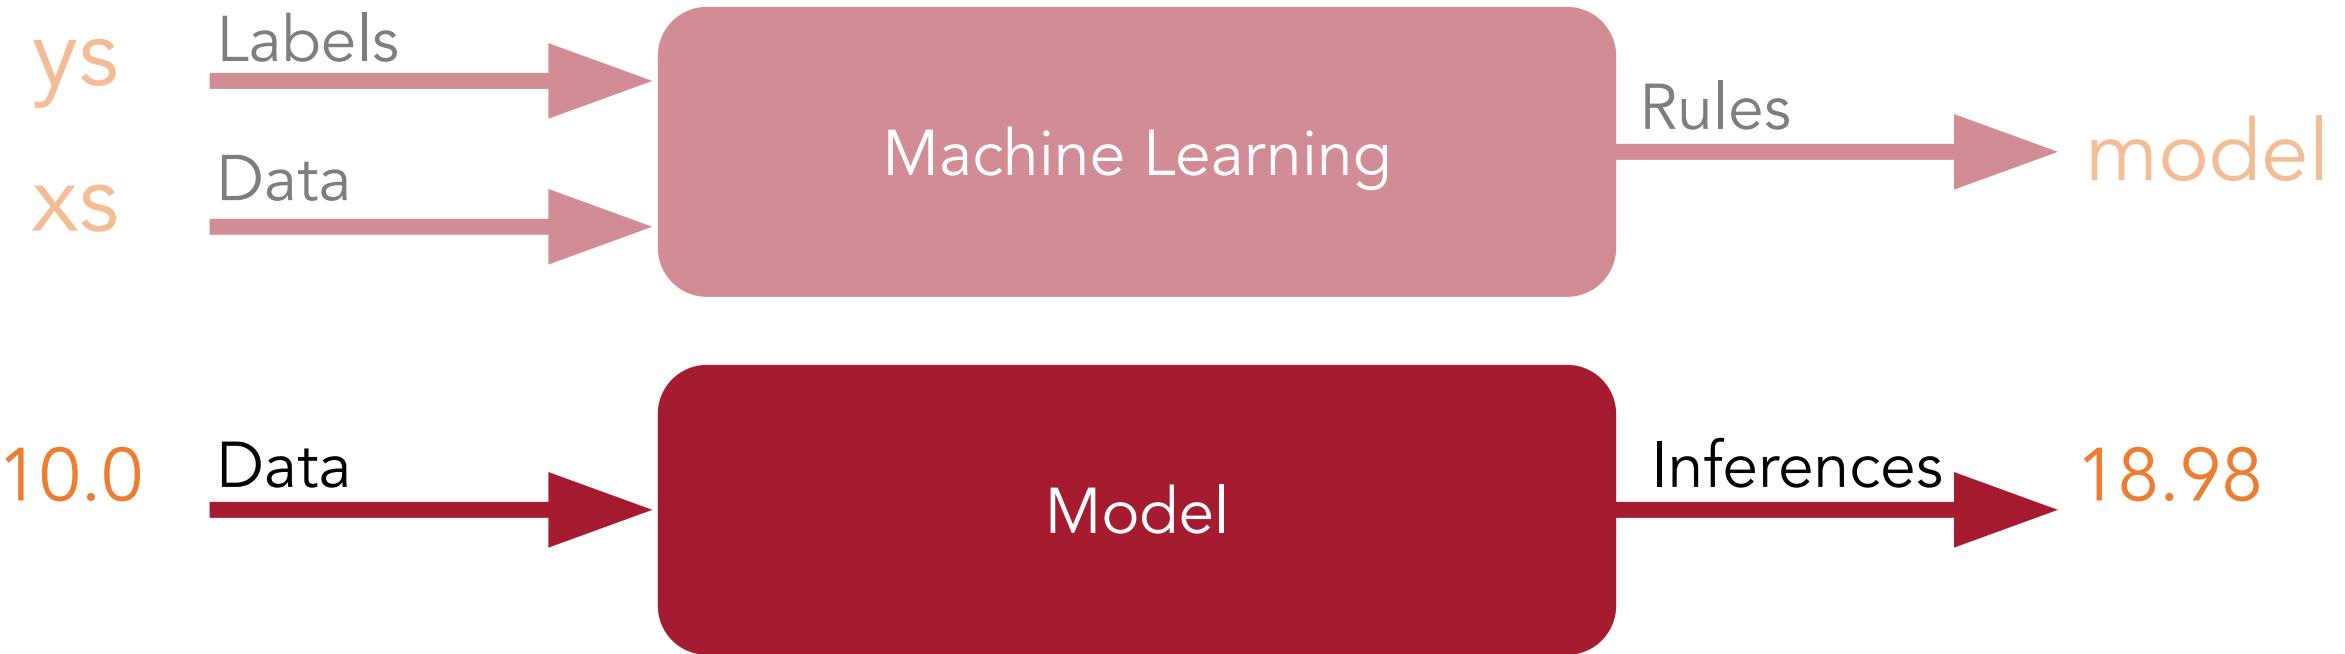
```
[[18.982813]]
```

```
[5] print(model.predict(xs))
```

```
[[-2.9897861]
 [-0.992277 ]
 [ 1.005232 ]
 [ 3.0027409]
 [ 5.00025  ]
 [ 6.997759 ]]
```

```
[6] print(my_layer.get_weights())
```

```
[array([[1.997509]], dtype=float32), array([-0.992277], dtype=float32)]
```



Linear Regression with a Single Neuron

colab.research.google.com
Regression.ipynb

```
[2] import tensorflow as tf
import numpy as np
from tensorflow import keras

# define a neural network with one neuron
# for more information on TF functions see: https://www.tensorflow.org/api\_docs
my_layer = keras.layers.Dense(units=1, input_shape=[1])
model = tf.keras.Sequential([my_layer])

# use stochastic gradient descent for optimization and
# the mean squared error loss function
model.compile(optimizer='sgd', loss='mean_squared_error')

# define some training data (xs as inputs and ys as outputs)
xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)

# fit the model to the data (aka train the model)
model.fit(xs, ys, epochs=500)
```

$$Y = 2X - 1$$

```
Epoch 500/500
1/1 [=====] - 0s 6ms/step - loss: 3.4704e-05
<keras.callbacks.History at 0x7f1d6cccd7f10>

[4] print(model.predict([10.0]))
[[18.982813]]

[5] print(model.predict(xs))
[[-2.9897861]
 [-0.992277]
 [ 1.005232]
 [ 3.0027409]
 [ 5.00025]
 [ 6.997759]]

[6] print(my_layer.get_weights())
[array([[1.997509]], dtype=float32), array([-0.992277], dtype=float32)]
```

$$Y = 1.9975X - 0.9922$$

Not perfect,
but good enough for most cases!

A Neuron in Tensor Flow

```
class Model(object):
    def __init__(self):
        self.w = tf.Variable(10.0)
        self.b = tf.Variable(10.0)

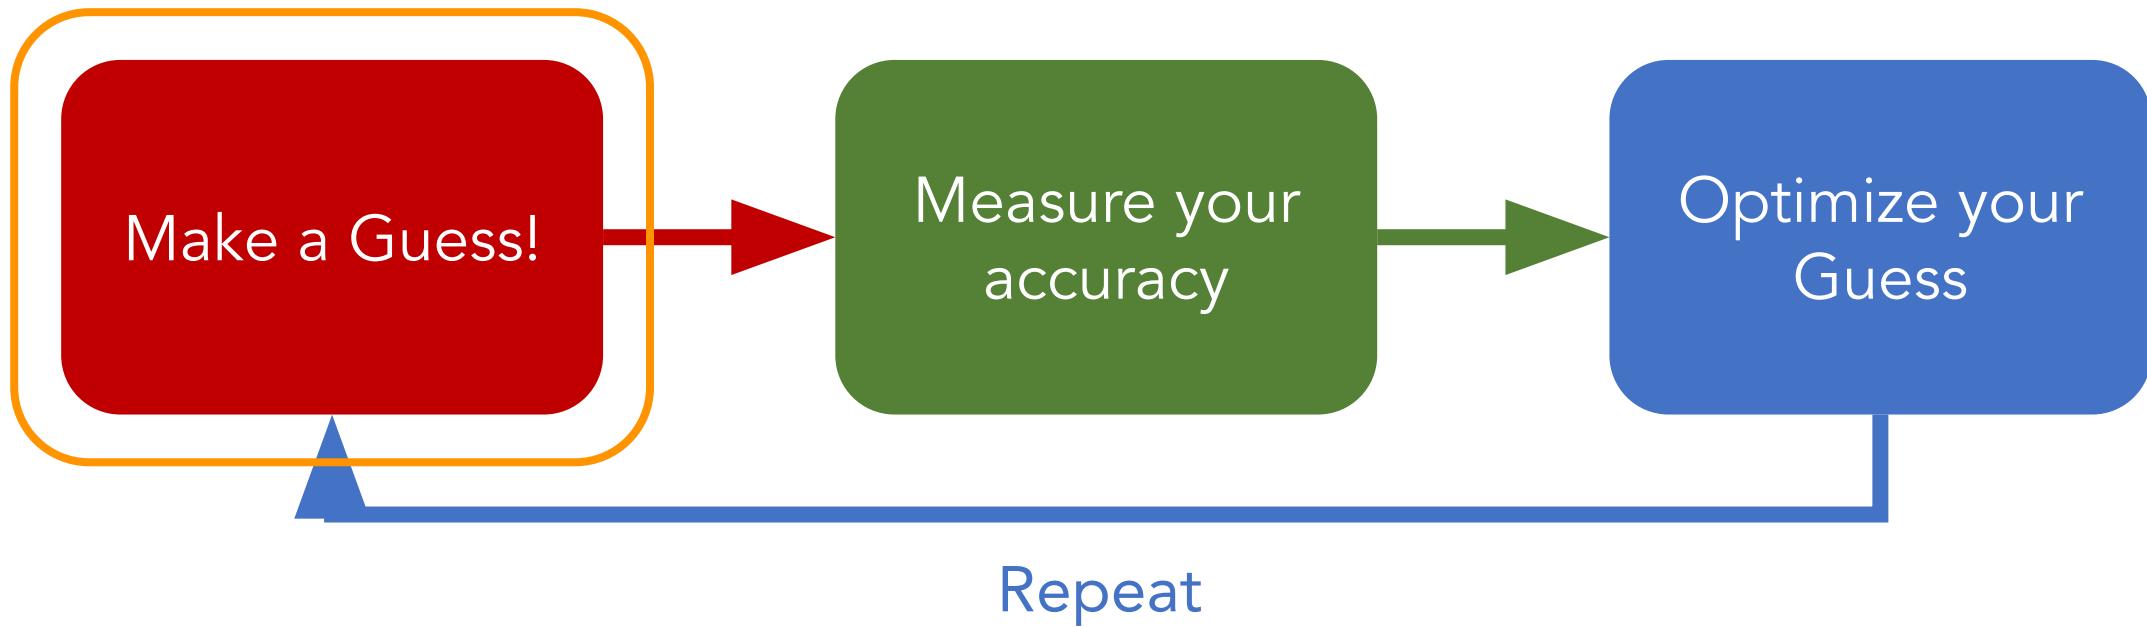
    def __call__(self, x):
        return self.w * x + self.b
```

A Neuron in Tensor Flow

```
class Model(object):
    def __init__(self):
        self.w = tf.Variable(10.0)
        self.b = tf.Variable(10.0)

    def __call__(self, x):
        return self.w * x + self.b
```

The Machine Learning Paradigm



A Neuron in Tensor Flow

```
model = Model()
xs = [-1.0, 0.0, 1.0, 2.0, 3.0, 4.0]
ys = [-3.0, -1.0, 1.0, 3.0, 5.0, 7.0]
print(model(xs))
```

```
Epoch 500/500
1/1 [=====] - 0s 6ms/step - loss: 3.4704e-05
<keras.callbacks.History at 0x7f1d6cccd7f10>
```

```
[4] print(model.predict([10.0]))
```



```
[[18.982813]]
```

```
[5] print(model.predict(xs))
```



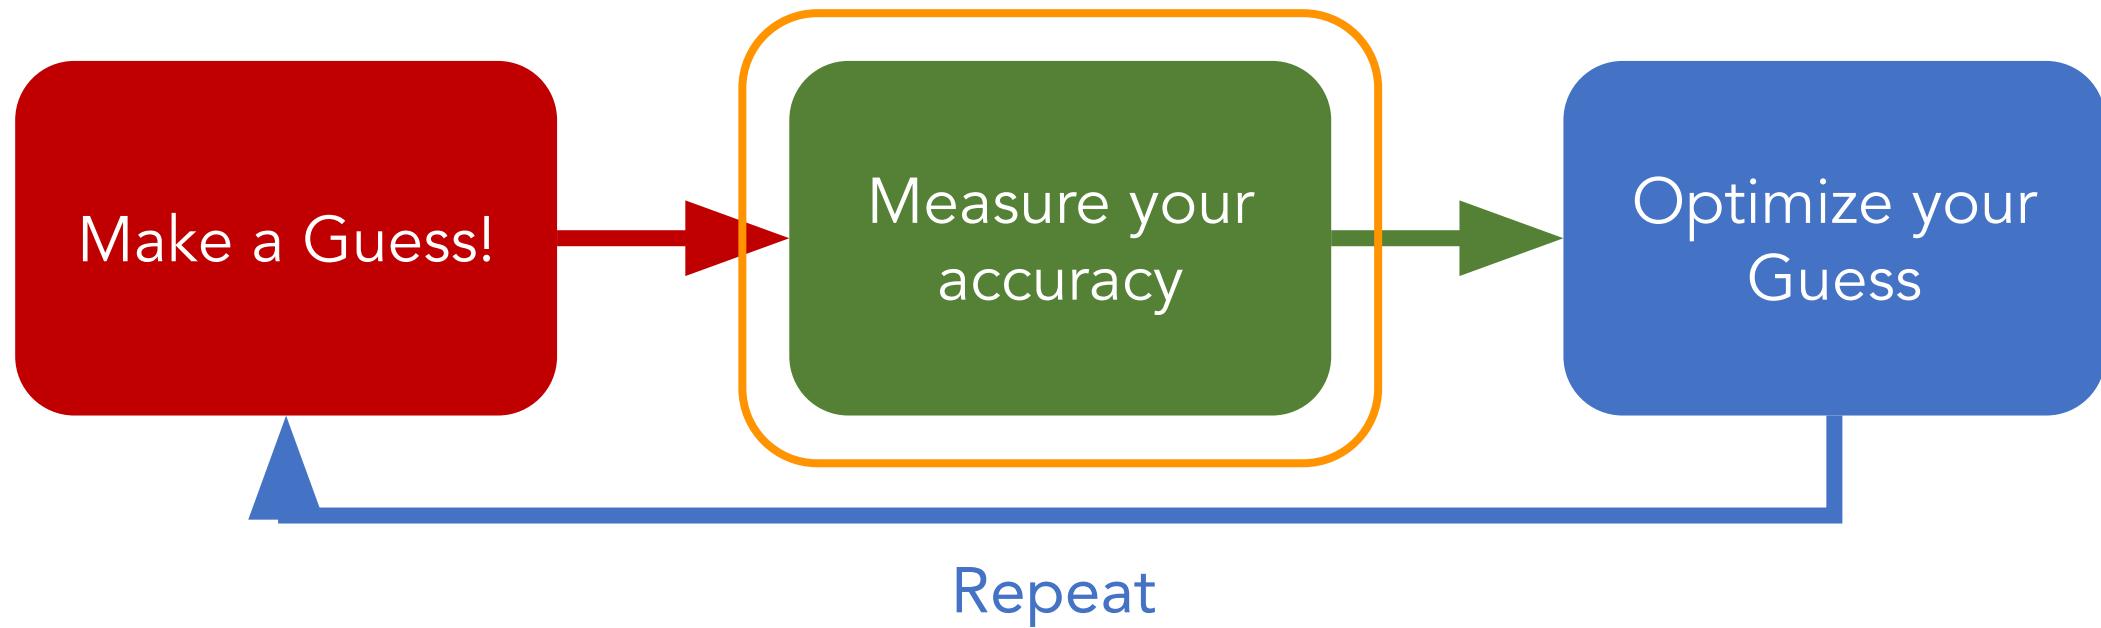
```
[[ -2.9897861]
 [ -0.992277 ]
 [  1.005232 ]
 [  3.0027409]
 [  5.00025  ]
 [  6.997759 ]]
```

```
[6] print(my_layer.get_weights())
```



```
[array([[1.997509]], dtype=float32), array([-0.992277], dtype=float32)]
```

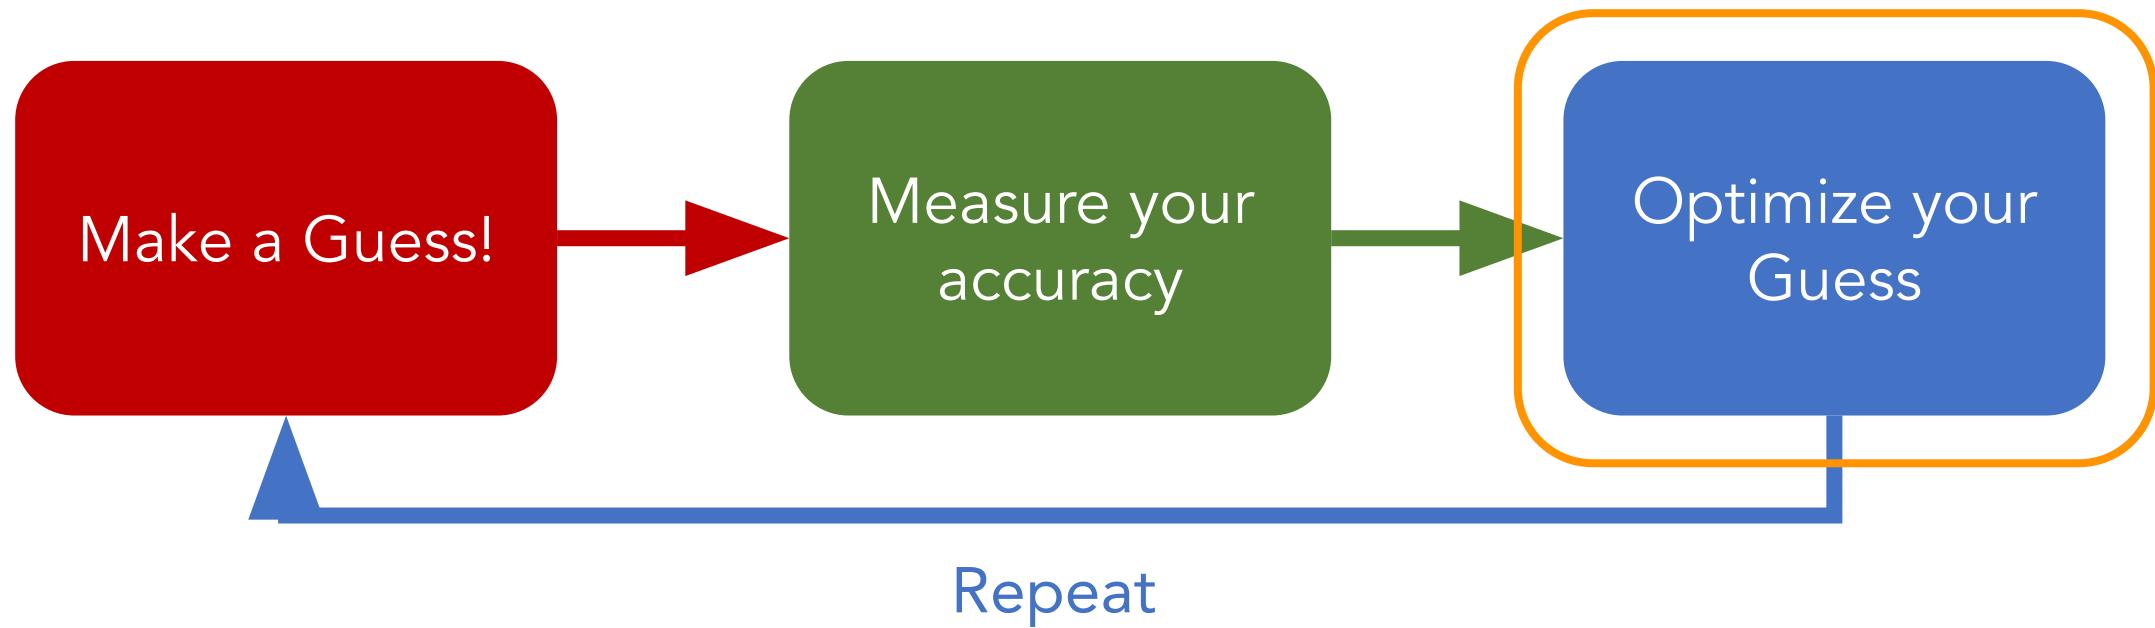
The Machine Learning Paradigm



```
def loss(predicted_y, target_y):  
    return tf.reduce_mean(tf.square(predicted_y - target_y))
```

```
def train(model, xs, ys, learning_rate):  
    with tf.GradientTape() as t:  
        current_loss = loss(model(xs), ys)  
  
    dw, db = t.gradient(current_loss, [model.w, model.b])  
    model.w.assign_sub(learning_rate * dw)  
    model.b.assign_sub(learning_rate * db)  
    return current_loss
```

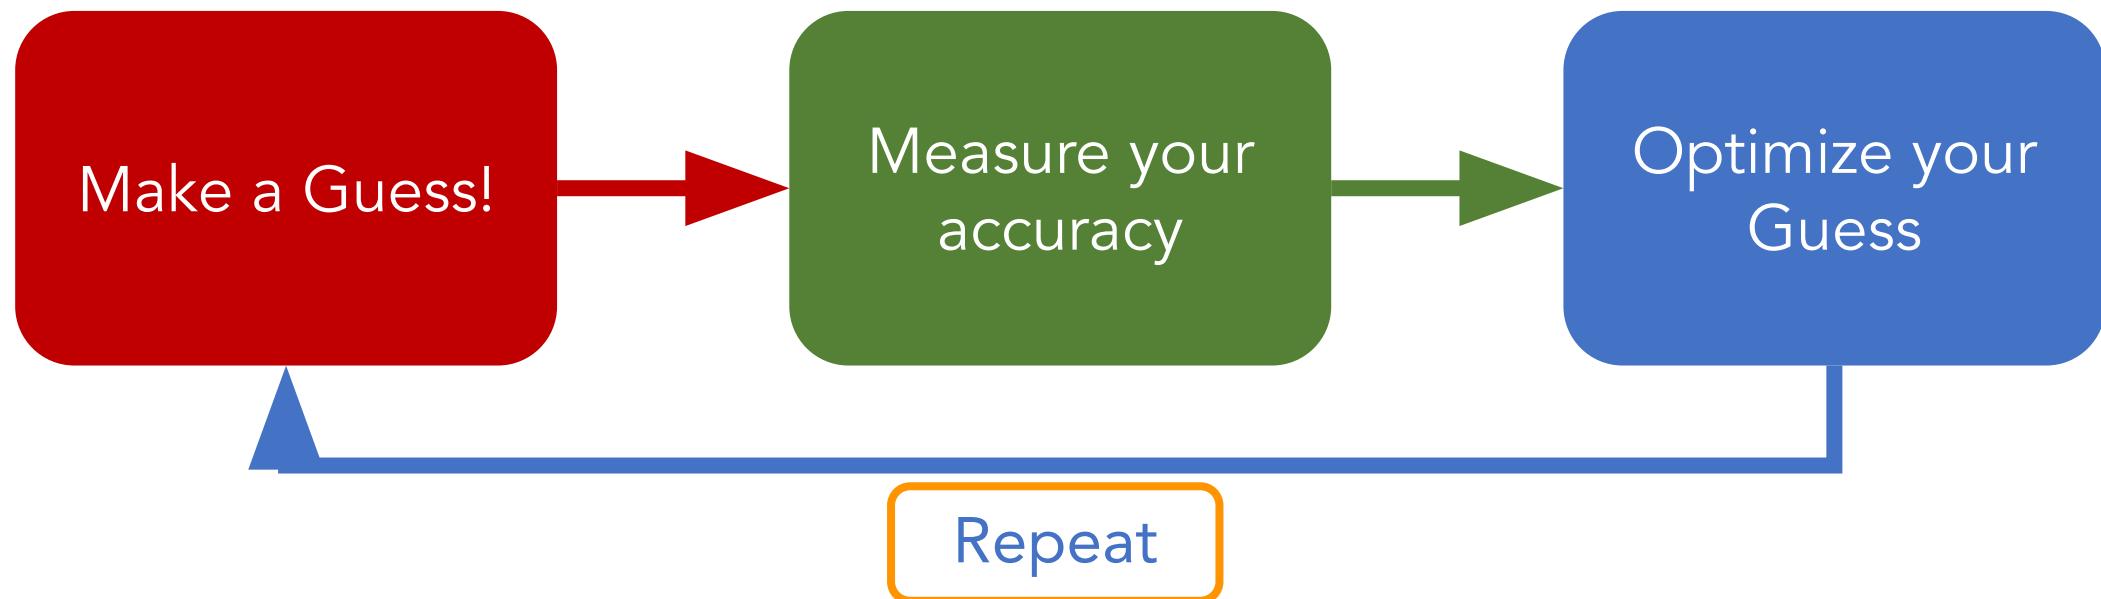
The Machine Learning Paradigm



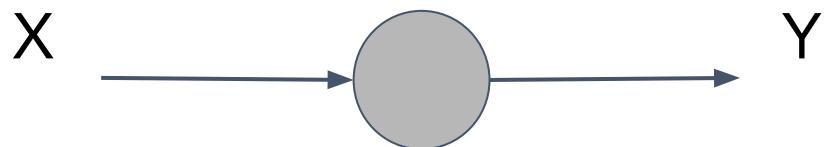
```
def train(model, xs, ys, learning_rate):  
    with tf.GradientTape() as t:  
        current_loss = loss(model(xs), ys)  
  
        dw, db = t.gradient(current_loss, [model.w, model.b])  
        model.w.assign_sub(learning_rate * dw)  
        model.b.assign_sub(learning_rate * db)  
    return current_loss
```

```
def train(model, xs, ys, learning_rate):  
    with tf.GradientTape() as t:  
        current_loss = loss(model(xs), ys)  
  
        dw, db = t.gradient(current_loss, [model.w, model.b])  
  
        model.w.assign_sub(learning_rate * dw)  
        model.b.assign_sub(learning_rate * db)  
  
    return current_loss
```

The Machine Learning Paradigm



```
for epoch in range(50):  
    current_loss = train(model, xs, ys, learning_rate=0.1)
```



$$y = f(x) = \mathbf{w}x + \mathbf{b}$$

$$y = 1.9975x - 0.9922$$



Gracias!

Prof. Diego Méndez Chaves, Ph.D

Associate Professor - Electronics Engineering Department
Director of the Master Program in Internet of Things
Director of the Master Program in Electronics Engineering
email: diego-mendez@javeriana.edu.co

