


Machine Learning Fundamentals

Workshop para América Latina y el Caribe (WALC)
Track 3 – Inteligencia Artificial Aplicada
November 12, 2024

Pontificia Universidad Javeriana – Bogotá, Colombia





School of Engineering

- 3.750 students (3000 undergrad)
 - 104 professors
 - 20.000 alumni
- 8 undergraduate programs
 - 7 specializations
 - 12 master programs
 - 2 doctoral programs



4

Departments



2

Centers of
Excellence

New Research and Laboratories Building

- 14.082 m² (700 for student workspaces)
- 93 meters high
- 15 floors and 3 basements
- 10.000 pieces of state-of-the-art equipment
- <https://ingenieria.javeriana.edu.co/nuestro-edificio>



1

Institutes



Diego Méndez Chaves, Ph.D

School of Engineering – Pontificia Universidad Javeriana

Website: <https://perfilesycapacidades.javeriana.edu.co/en/persons/diego-mendez-chaves>

email: diego-mendez@javeriana.edu.co

- Associate Professor at the Department of Electronics Engineering.
- Director of the Master Program in Internet of Things.
- Director of the Master Program in Electronics Engineering.
- Technical Director of the Center of Excellence and Adoption in IoT (CEA-IoT)
- Research associate at the Marconi Lab in the International Centre for Theoretical Physics (ICTP), Trieste - Italy.
- TinyML Academic Network – Coordinator.
- Research interests: IoT, embedded systems, wireless sensor networks, participatory sensing, digital systems design and embedded operating systems.

Track3

Inteligencia
Artificial Aplicada



Agenda

- Parte 1:
 - Introducción
 - El paradigma de ML
 - Exploración de la función de pérdida y costo
 - Redes neuronales artificiales
- Parte 2:
 - Dense Neural Networks – Regresión
 - Dense Neural Networks – Clasificación
 - Métricas de ML





Alexa, play some
rock music!

Playing the Rock
Hits playlist

AI vs. ML vs. DL

Artificial Intelligence



Machine Learning



Deep Learning




Artificial intelligence (AI): any technique that enables computers to mimic human intelligence.


Machine learning (ML): a subset of AI that uses techniques that enable machines to use experience to improve at tasks.

Deep learning (DL): a subset of ML based on artificial neural networks (ANN). The learning process is deep because of its structure.

General Steps for Machine Learning

On a high level, the craft of creating machine learning (ML) processes is comprised of several steps:





“The future of ML is *tiny* and bright.”

3 main components

“**Edge AI** is a truly complete technology. As a topic, it makes use of knowledge from everything from the physical properties of semiconductor electronics all the way up to the engineering of high-level architectures that span devices and the cloud. It demands expertise in the most cutting-edge approaches to **artificial intelligence and machine learning** along with the most venerable skills of bare-metal **embedded software engineering**. It makes use of the entire history of computer science and electrical engineering, laid out end to end.”




Situnayake, Daniel; Plunkett, Jenny
AI at the Edge (pp. 215-216)
O'Reilly Media

3 main components



Interactions

Given your understanding of things at these various intersections, you will have a deep understanding for **how to apply TinyML**



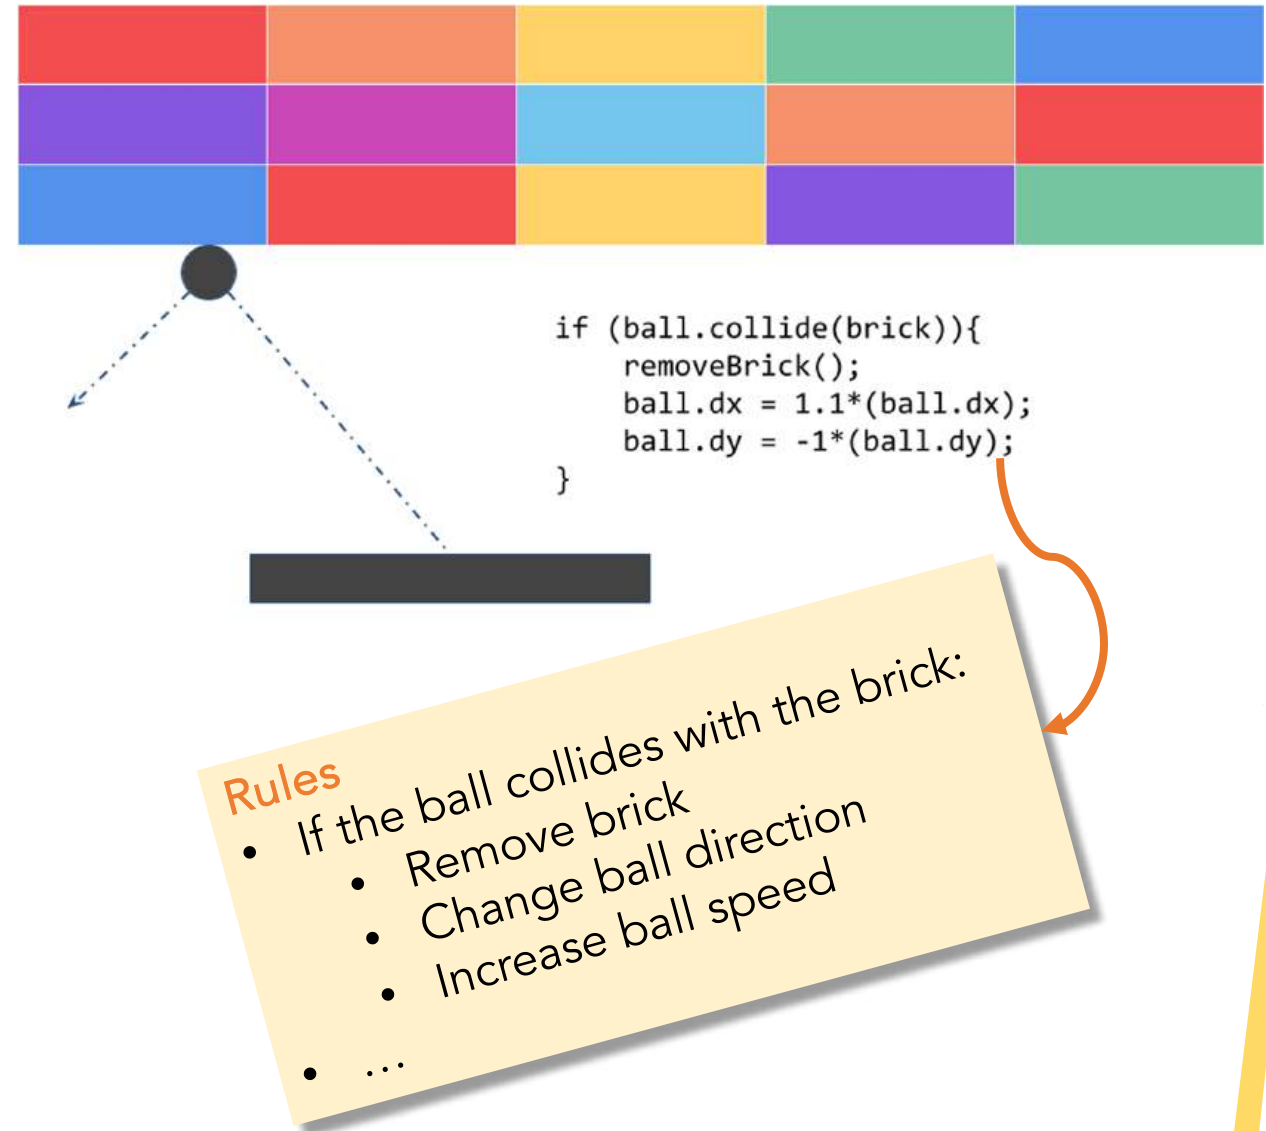


The Machine Learning Paradigm



Explicit Coding

- **Defining rules** that determine behavior of a program
- Everything is **pre-calculated and pre-determined** by the programmer
- **Scenarios are limited** by program complexity



The Traditional Programming Paradigm



Consider Activity Detection



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```




```
// ???
```

Way too
complex
to code!

The Traditional Programming Paradigm



The Machine Learning Paradigm



Activity Detection with Machine Learning



```
0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010
```

Label = WALKING



```
1010100101001010101
0101010010010010001
0010011111010101111
1010100100111101011
```

Label = RUNNING



```
1001010011111010101
1101010111010101110
1010101111010101011
1111110001111010101
```

Label = BIKING



```
1111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110
```

Label = GOLFING

The Machine Learning Paradigm



```
0101001010100101010
1001010101001011101
0100101010010101001
0101001010100101010
```

Label = WALKING



```
1010100101001010101
0101010010010010001
0010011111010101111
1010100100111101011
```

Label = RUNNING



```
1001010011111010101
1101010111010101110
1010101111010101011
1111110001111010101
```

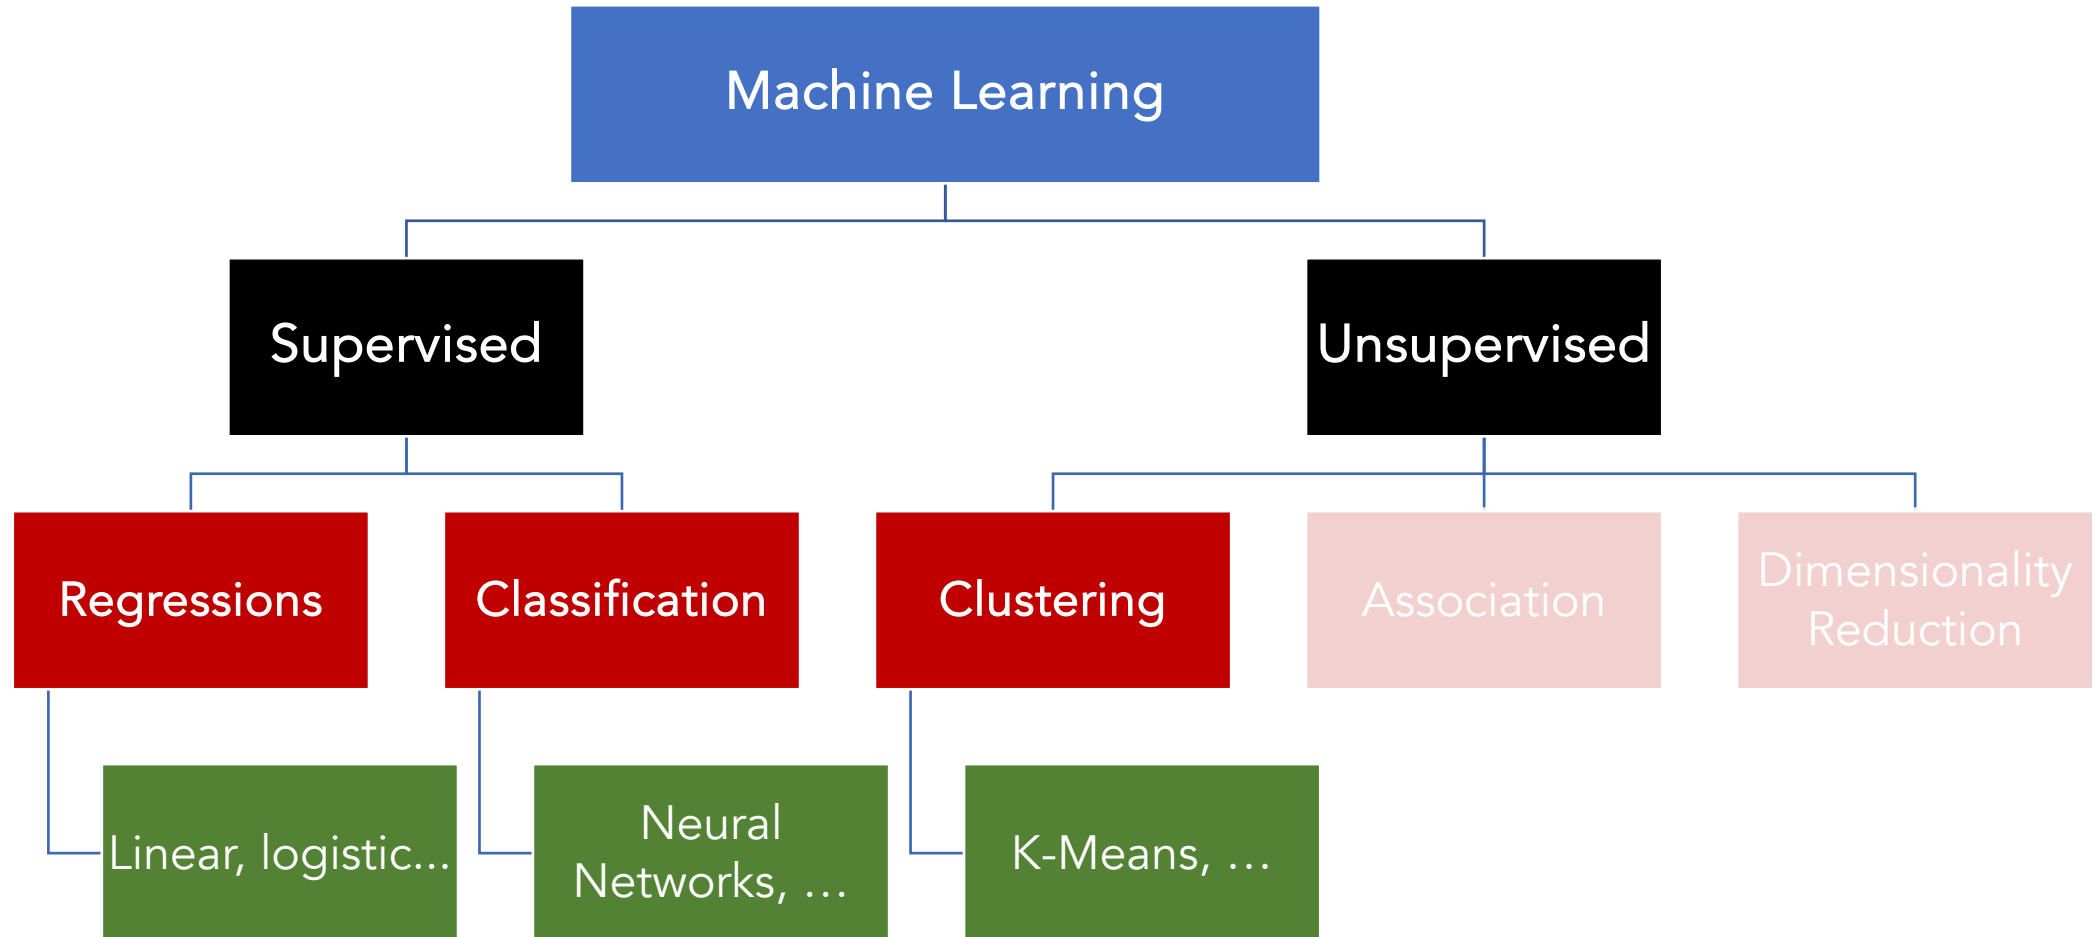
Label = BIKING



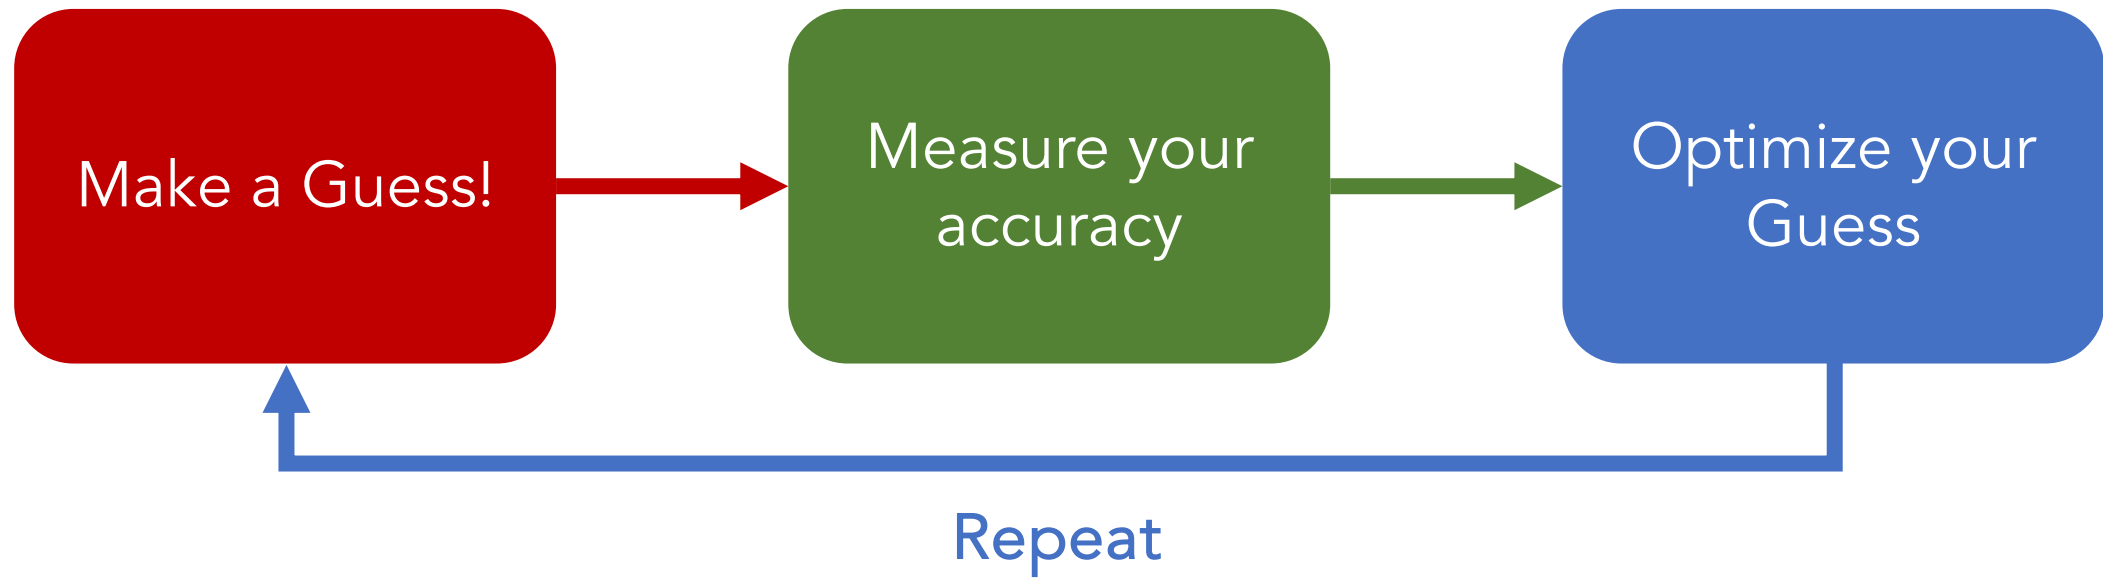
```
1111111111010011101
0011111010111110101
0101110101010101110
1010101010100111110
```

Label = GOLFING

Two Approaches




The Machine Learning Paradigm



The Machine Learning Paradigm



The Machine Learning Paradigm



How good is your model?

a way to measure your accuracy

Matching X to Y

$$X = \{-1, 0, 1, 2, 3, 4\}$$

$$Y = \{-3, -1, 1, 3, 5, 7\}$$



Make a guess!

$$Y = 3X - 1$$

$$X = \{-1, 0, 1, 2, 3, 4\}$$

$$\text{My } Y = \{-4, -1, 2, 5, 8, 11\}$$

How good is the guess?


$$Y = 3X - 1$$

$$X = \{-1, 0, 1, 2, 3, 4\}$$


$$\text{My } Y = \{-4, -1, 2, 5, 8, 11\}$$

$$\text{Real } Y = \{-3, -1, 1, 3, 5, 7\}$$


Let's measure it!




Let's measure it!




Let's measure it!





Let's measure it!



Houston, we have a
problem!



What if we square² them?




Total that (Σ) and take
the square root $\sqrt{}$

$$\text{sqrt}(1 + 1 + 4 + 9 + 16)$$

$$= \text{sqrt}(31)$$

$$= 5.57$$



Make another guess!

$$Y = 2X - 2$$

$$X = \{-1, 0, 1, 2, 3, 4\}$$

$$\text{My } Y = \{-4, -2, 0, 2, 4, 6\}$$

$$\text{Real } Y = \{-3, -1, 1, 3, 5, 7\}$$

$$\text{Diff}^2 = \{1, 1, 1, 1, 1\}$$



Get the same
difference, repeat the
same process.

$$\text{sqrt}(1 + 1 + 1 + 1 + 1)$$

$$= \text{sqrt}(5)$$

$$= 2.23$$



Make another guess!

$$Y = 2X - 1$$

$$X = \{-1, 0, 1, 2, 3, 4\}$$

$$\text{My } Y = \{-3, -1, 1, 3, 5, 7\}$$

$$\text{Real } Y = \{-3, -1, 1, 3, 5, 7\}$$

$$\text{Diff}^2 = \{0, 0, 0, 0, 0\}$$



Make another guess!

$$Y = 2X - 1$$

$$X = \{-1, 0, 1, 2, 3, 4\}$$


$$\text{My } Y = \{-3, -1, 1, 3, 5, 7\}$$

$$\text{Real } Y = \{-3, -1, 1, 3, 5, 7\}$$

$$\text{Diff}^2 = \{0, 0, 0, 0, 0\}$$

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

Goal is to Minimize MSE (Mean Squared Error)



Finding out the best solution

Trial and error approach

Loss
Function


Parameter



The image shows a graph of a loss function. A blue U-shaped curve represents the loss, with its minimum point located on a vertical dotted line. A horizontal dotted line serves as the x-axis, which is labeled 'Parameter'. A red double-headed arrow is positioned above the x-axis, centered under the minimum of the curve, indicating the range of parameters being considered. The text 'Loss Function' is positioned to the left of the curve, and 'Parameter' is positioned to the right of the x-axis label.

Loss
Function

Minimum of
Loss Function





Loss
Function




Loss
Function

Gradient of
value



Loss
Function

Move in Direction of Gradient
Learning Rate is size of the step to take




Loss
Function

End up here




Loss
Function

Get the
gradient



Loss
Function

Move in Direction of Gradient




Loss
Function


End Up here




Loss
Function




Loss
Function




Loss
Function




Loss
Function




Loss
Function




Loss
Function




Loss
Function



Loss
Function



Loss
Function



Loss
Function

Move in Direction of Gradient



Loss
Function

Move in Direction of Gradient



Loss
Function

Move in Direction of Gradient



Loss
Function

Move in Direction of Gradient




Loss
Function

Move in Direction of Gradient




It is important to choose the correct Learning Rate (size of the step)



If the **Learning Rate** is too small it may take a long time to reach the minimum




If the **Learning Rate** is too large we may never reach the minimum



Gradient Descent algorithm





Gradient Descent for Two Parameters




A single minima
Global minima

Gradient Descent for Two Parameters






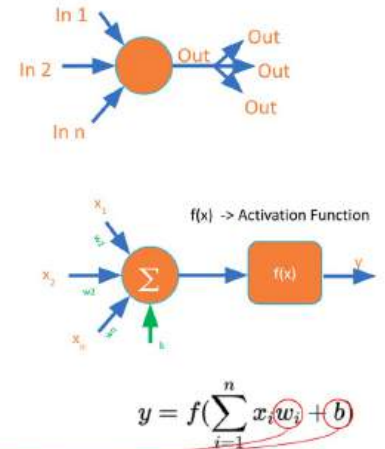

Artificial Neural Networks



What is an Artificial Neural Network (ANN)?



Neuron (Perceptron)



Parameters


$$y = f\left(\sum_{i=1}^n x_i w_i + b\right)$$

Dense Neural Network
(DNN)

Sequential

A neuron

a neuron's output is a function of its inputs (in this case only one)



$$y = f(x) = wx + b$$

There are only **two parameters** to adjust:
The **weight** for each input and a **bias**

First scenario: a regression

Linear Regression with a Single Neuron

colab.research.google.com

Regression.ipynb

```
[2] import tensorflow as tf
import numpy as np
from tensorflow import keras
```

```
# define a neural network with one neuron
# for more information on TF functions see: https://www.tensorflow.org/api\_docs
my_layer = keras.layers.Dense(units=1, input_shape=[1])
model = tf.keras.Sequential([my_layer])
```

```
# use stochastic gradient descent for optimization and
# the mean squared error loss function
model.compile(optimizer='sgd', loss='mean_squared_error')
```

```
# define some training data (xs as inputs and ys as outputs)
xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
```

```
# fit the model to the data (aka train the model)
model.fit(xs, ys, epochs=500)
```

1 layer, 1 neuron, 1 input

Stochastic gradient descent

Inputs and outputs (labels)

Train the model



Open in Colab



Linear Regression with a Single Neuron

colab.research.google.com

Regression.ipynb

```
[2] import tensorflow as tf
import numpy as np
from tensorflow import keras

# define a neural network with one neuron
# for more information on TF functions see: https://www.tensorflow.org/api\_docs
my_layer = keras.layers.Dense(units=1, input_shape=[1])
model = tf.keras.Sequential([my_layer])

# use stochastic gradient descent for optimization and
# the mean squared error loss function
model.compile(optimizer='sgd', loss='mean_squared_error')

# define some training data (xs as inputs and ys as outputs)
xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)

# fit the model to the data (aka train the model)
model.fit(xs, ys, epochs=500)
```

```
Epoch 500/500
1/1 [=====] - 0s 6ms/step - loss: 3.4704e-05
<keras.callbacks.History at 0x7f1d6ccd7f10>
```

```
[4] print(model.predict([10.0]))

[[18.982813]]
```

```
[5] print(model.predict(xs))

[[-2.9897861]
 [-0.992277 ]
 [ 1.005232 ]
 [ 3.0027409]
 [ 5.00025  ]
 [ 6.997759 ]]
```

```
[6] print(my_layer.get_weights())

[array([[1.997509]], dtype=float32), array([-0.992277], dtype=float32)]
```



Linear Regression with a Single Neuron

colab.research.google.com

Regression.ipynb

```
[2] import tensorflow as tf
import numpy as np
from tensorflow import keras

# define a neural network with one neuron
# for more information on TF functions see: https://www.tensorflow.org/api\_docs
my_layer = keras.layers.Dense(units=1, input_shape=[1])
model = tf.keras.Sequential([my_layer])

# use stochastic gradient descent for optimization and
# the mean squared error loss function
model.compile(optimizer='sgd', loss='mean_squared_error')

# define some training data (xs as inputs and ys as outputs)
xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)

# fit the model to the data (aka train the model)
model.fit(xs, ys, epochs=500)
```

$$Y = 2X - 1$$

```
Epoch 500/500
1/1 [=====] - 0s 6ms/step - loss: 3.4704e-05
<keras.callbacks.History at 0x7f1d6ccd7f10>
```

```
[4] print(model.predict([10.0]))

[[18.982813]]
```

```
[5] print(model.predict(xs))

[[-2.9897861]
 [-0.992277 ]
 [ 1.005232 ]
 [ 3.0027409]
 [ 5.00025 ]
 [ 6.997759 ]]
```


```
[6] print(my_layer.get_weights())


[array([[1.997509]], dtype=float32), array([-0.992277], dtype=float32)]
```

$$Y = 1.9975X - 0.9922$$

Not perfect,
but good enough for most cases!

The Machine Learning Paradigm





$$y = f(x) = wx + b$$

$$y = 1.9975x - 0.9922$$





Gracias!

Prof. Diego Méndez Chaves, Ph.D

Associate Professor - Electronics Engineering Department
Director of the Master Program in Internet of Things
Director of the Master Program in Electronics Engineering

<https://perfilesycapacidades.javeriana.edu.co/en/persons/diego-mendez-chaves>

diego-mendez@javeriana.edu.co